Programming in Oberon

A derivative of Programming in Modula-2 (1982)

Niklaus Wirth

Table of Contents

Preface

This text is an introduction to programming in general, and a guide to programming in the language Oberon in particular. It is primarily oriented towards people who have already acquired some basic knowledge of programming and would like to deepen their understanding in a more structured way. Nevertheless, an introductory chapter is added for the benefit of the beginner, displaying in a concise form some of the fundamental concepts of computers and of programming them. The text is therefore also suitable as a self-contained tutorial. The notation used is Oberon, which lends itself well for a structured approach and leads the student to a working style that has generally become known under the heading of *structured programming*.

As a manual for programming in Oberon, the text covers (almost) all facilities of this language. Part 1 covers the basic notions of variable, expression, assignment, conditional and repeated statement, and the data structure of the indexable array. Together with Part 2, which introduces the important concept of the procedure or subroutine, it contains essentially the material commonly discussed in introductory programming courses. Part 3 concerns data types and structures and constitutes the core of an advanced course on programming. Part 4 introduces the notion of a module, a concept that is fundamental in the design of large programming systems, and to programming in teams. The two basic notions of files and texts are introduced in the forms of modules containing the operations performed on files and texts. Texts are presented as our standard medium for program input and output.

The language Oberon, published in 1988, has a long history. It is a descendant of Algol 60 (1960), Pascal (1970), and Modula (1979). Algol 60 [1] was designed by an international committee of 13 scientists, and it was the first language to be specified with a rigorous formalism, a *syntax*, in machine-independent form. Algol was largely oriented toward numerical algorithms, as were computers at that time. Pascal [2] was the result of an enduring debate about widening the range of application of Algol, and it became widely used, in particular but not only, in education. Modula-2 [3] introduced the concept of modules, the parts of large systems connected by clearly specified interfaces. The module allows to hide details of procedures and variables from its users (clients), and it embodies the principle of *information hiding*. The language Oberon [4] emerged from the urge to reduce the complexity of programming languages, of Modula in particular. This effort resulted in a remarkably concise language. The extent of Oberon, the number of its features and constructs, is smaller even than that of Pascal. Yet it is considerably more powerful.

The one feature that was added in Oberon was the extensibility of data types (record types). Whereas in strongly types languages, such as Algol, Pascal, and Modula, every constant, variable, or function has a fixed type, recognizable from the program text, Oberon allows to define hierarchies of types, and to determine the actual type of a variable (within the hierarchy) at run-time. This, together with records containing fields of procedural types, is the stem of *object-oriented programming*. Such records are then called *objects*, and the procedural fields are called *methods*. Here, we do not touch upon this subject.

Zürich, 1. October 2004   N.W.

1. P. Naur, Ed. Revised Report on the Algorithmic Language Algol 60. *Comp. J. 5*, 349-367 (1962), and *Comm. ACM, 6* (1963) 1 – 17.

2. N. Wirth. The Programming Language Pascal. *Acta Informatica, 1* (1971), 35 – 63.

3. N. Wirth. *Programming in Modula-2*. Springer-Verlag, 1982. ISBN 0-387-50150-9

4. N. Wirth. The Programming Language Oberon. *Software - Practice and Experience, 18*, 7, (July 1988), 671- 690.

Part 1

1. Introduction

Although this manual assumes that its reader is already familiar with the basic notions of computer and programming, it may be appropriate to start out with the explanation of some concepts and their terminology. We recognize that - with rare exceptions - programs are written - more appropriately: designed - with the purpose of being interpreted by a computer. The computer then performs a process, i.e. a sequence of actions, according to the specifications given by that program. The process is also called a *computation*.

The program itself is a *text*. Since it specifies a usually fairly complex process, and must do so with utmost precision and care for all details, the meaning of this text must be specified very precisely. Such precision requires an exact formalism. This formalism has become known as a *language*. We adopt this name, although a language is normally spoken and much less precisely defined. Our purpose here is to learn the formalism or language *Oberon*. It has a long tradition among programming languages. Its ancestor was Algol 60, followed by Pascal and Modula-2.

A program usually specifies a process that causes its interpreter, i.e. the computer, to read data (the so-called *input*) from some sources and to vary its subsequent actions according to the accepted data. This implies that a program does not only specify a (single) process, but an entire - usually unbounded - class of computations. We have to ensure that these processes act according to the given specifications (or should we say expectations?) in all cases of this class. Whereas we could verify that this specification is met in the case of a single computation, this is impossible in the general case, because the class of all permitted processes is much too large. The conscientious programmer ensures the correctness of his program by careful design and analysis. Careful design is the essence of professional programming.

The task of designing a program is further complicated by the fact that the program not only must describe an entire class of computations, but often should also be interpreted (executed) by different interpreters (computers). At earlier times, this required the manual transcription of the program from its source form into different computer codes, taking into account their various characteristics and limitations. The difficulties have been drastically reduced, albeit not eliminated, by the creation of high level languages with formal definitions and the construction of automatic translators converting the program into the codes of the various computers.

In principle, the formal language should be defined in an abstract, perhaps axiomatic fashion without reference to an actual computer or interpretation mechanism. If this were achieved, the programmer would have to understand the formal language only. However, such generality is costly and often restrictive, and in many cases the programmer should still know the principal characteristics of his computer(s). Nevertheless, the qualified programmer will make as little reference to specific computer characteristics as possible and rely exclusively on the rules of the formal language in order to keep his program general and portable. The language Oberon assists in this task by confining computer dependencies to specific objects, by allowing to encapsulate them in specific, small parts of a program text.

From the foregoing it follows that a translation process lies between the program's formulation and its interpretation. This process is called a *compilation*, because it condenses the program's source text into a cryptic computer code. The quality of this compilation may be crucial to the efficiency of the program's ultimate interpretation. We stress the fact that there may be many compilers for a given language (even for the same computer). Some may be more efficient than others. We recognize that efficiency is a characteristic of implementations rather than the language. It therefore is important to distinguish between the concepts of language and implementation.

We summarize:
- A program is a piece of *text*.
- The program specifies *computations* or processes.
- A process is performed by an interpreter, usually a *computer*, interpreting (executing) the program.
- The meaning of the program is specified by a formalism called *programming language*.
- A program specifies a class of computations, the *input data* acting as parameter of each individual process.

-    Prior to its execution, a program text is translated into computer code by a compiler. This process is called a *compilation*.

Program design includes ensuring that all members of this class of computations act according to specification. This is done by careful analytic *verification* and by selective empirical testing of characteristic cases.

Programs should refrain from making reference to characteristics of specific interpreters (computers) whenever possible. Only the lack of such reference ensures that their meaning can be derived from rules of the language.

A compiler is a program translating programs from their source form to specific computer codes. Programs need to be compiled before they are executed. Programming in the wider sense not only includes the formulation of the program, but also the concrete preparation of the text, its compilation, correction of errors, so-called *debugging*, and the planning of tests. The modern programmer uses many tools for these tasks, including text editors, compilers, and debuggers. He also has to be familiar with the environment of these components. We shall not describe these aspects, but concentrate on the *language Oberon*.

2. A First Example

Let us follow the steps of development of a simple program and thereby explain some of the fundamental concepts of programming and of the basic facilities of Oberon. The task shall be, given two natural numbers x and y, to compute their greatest common divisor (gcd). The mathematical knowledge needed for this problem is the following:

1. if x equals y, x (or y) is the desired result
2. the gcd of two numbers remains unchanged, if we replace the larger number by the difference of the numbers, i.e. subtract the smaller number form the larger one.

Expressed in mathematical terms, these rules take the form

1. $gcd(x, x) = x$
2.   if $x > y$, $gcd(x, y) = gcd(x-y, y)$

The basic recipe, the so-called *algorithm*, is then the following: Change the numbers x and y according to rule 2 such that their difference decreases. Repeat this until they are equal. Rule 2 guarantees that the changes are such that gcd(x,y) always remains the same, and rule 1 guarantees that we finally find the result.
Now we must put these recommendations into terms of Oberon. A first attempt leads to the following sketch. Note that the symbol # means "unequal".

```
WHILE x # y DO
  "apply rule 2, reducing the difference and maintaining x > 0 and y > 0"
END
```

The sentence within quotes is plain English. The second version refines the first version by replacing the English by formal terms:

```
WHILE x # y DO
  IF x > y THEN x := x-y ELSE y := y-x END
END
```

This piece of text is not yet a complete program, but it shows already the essential characteristic of a structured programming language. Version 1 is a statement, and this statement contains another, subordinate statement (within quotes). In version 2 this is elaborated, and yet further subordinate statements emerge (expressing the replacement of a value x by another value x-y). This hierarchy of statements expresses the underlying structure of the algorithm. It becomes explicit due to the structure of the language, allowing the nesting of components of a program. It is therefore important to know the language's structure (syntax) in full detail. Textually we express nesting or subordination by appropriate indentation. Although this is not required by the rules of the language, it helps in the understanding of a text very considerably.

Reflecting an algorithm's inherent structure by the textual structure of the program is a key idea of structured programming. It is virtually impossible to recognise the meaning of a program when its structure is removed, such as done by a compiler when producing computer code. And we should keep

in mind that a program is worthless, unless it exists in some form in which a human can understand it and gain confidence in its design.

We now proceed towards the goal of producing a complete program from the above fragment. We realize that we need to specify an action that assigns initial values to the variables x and y, as well as an action that makes the result visible. For this purpose we should actually know about a computer's facilities to communicate with its user. Since we do not wish to refer to a specific machinery, and particularly not in such a frequent and important case as the generation of output, we introduce abstractions of such communication facilities. These facilities are not directly part of the language, but are procedures declared in some (library) modules, to which all programs have access. We postulate read and write procedures as shown in the following example, and will assume that data are thus read from a keyboard and written on a display.

```
Texts.Scan(S); x := S.i;
Texts.Scan(S); y := S.i;
WHILE x # y DO
  IF x > y THEN x := x-y ELSE y := y-x END
END;
Texts.WriteInt(W, x, 6)
```

The procedure *Scan* scans an input text and reads a (non-negative) integer (S.i). The procedure *WriteInt* outputs an integer as specified by its first parameter (x). The second parameter (6) indicates the number of digits available for the representation of this value in the output text. Both procedures are taken from the imported (library) module *Texts*, and they use a locally declared scanner S and a globally declared writer W, connecting to an input source (the text following the command) and output sink (the Log text) imported from the environment *Oberon*. Details are explained in Chapter 22 of this tutorial.

In the next and final version we complete our text such that it becomes a genuine Oberon text:

```
PROCEDURE Gcd*;
  VAR x, y: INTEGER; S: Texts.Scanner;
BEGIN Texts.OpenScanner(S, Oberon.Par.text, Oberon.Par.pos);
  Texts.Scan(S); x := S.i; Texts.WriteString(W, "  x ="); Texts.WriteInt(W, x, 6);
  Texts.Scan(S); y := S.i; Texts.WriteString(W, "  y ="); Texts.WriteInt(W, y, 6);
  WHILE x # y DO
    IF x > y THEN x := x-y ELSE y := y-x END
  END;
  Texts.WriteString(W, " gcd ="); Texts.WriteInt(W, x, 6); Texts.WriteLn(W);
  Texts.Append(Oberon.Log, W.buf)
END Gcd.
```

The essential additions in this step are *declarations*. In Oberon, all names of objects occurring in a program, such as variables and constants, have to be declared. A declaration introduces the object's identifier (name), specifies the kind of the object (whether it is a variable, a constant, or something else) and indicates general, invariant properties, such as the type of a variable or the value of a constant.

The entire piece of text is called a *procedure*, given a name, and has the following format (see also Ch. 10):

```
PROCEDURE name;
  <declarations>
BEGIN
  <statements>
END name
```

In this and the following examples, a particular property of the Oberon system becomes apparent. Whereas in older languages, the notions of program and executable unit were identical, we distinguish carefully between them in Oberon terminology: What used to be called a *program* is now called a *module*, a system unit typically incorporating variables and executable statements resident in a system (see also Ch. 19). Compilers accept a module as a *compilable unit* of text. An *executable unit* of program we call a *procedure*. In the Oberon system, any parameterless procedure can be used as a *command*, if its name is marked with an asterisk. It can be activated by clicking on its name visible anywhere on the display. A module may contain one or several commands. As a consequence, the

preceding and following examples of "programs" appear in the form of procedures (commands), and we will always assume that they are embedded in a module importing two service modules *Texts* and *Oberon* and containing the declaration of a writer *W* for output:

```
MODULE Pattern;
  IMPORT Texts, Oberon;
  VAR W: Texts.Writer;

  PROCEDURE command*;  (*such as Gcd*)
  BEGIN …  (*using W for output*) …
    Texts.Append(Oberon.Log, W.buf)
  END command;
BEGIN Texts,OpenWriter(W)
END Pattern.
```

A few more comments concerning our *Gcd* example are in order. As already mentioned, the procedures *WriteLn, WriteString, WriteInt* and *Scan* are not part of the language Oberon itself. They are defined in a module called *Texts* which is presumed to be available. The often encountered modules *Texts*, *Files* and *Oberon* will be defined and explained in Chapter 22 of this book. Here we merely point out that they need to be imported in order to be known in a program. This is done by including the names of these modules in the import list in the importing module's heading.

The procedure *WriteString* outputs a string, i.e. a sequence of characters (enclosed in quotes). The procedure *WriteLn* inserts a line break in the output text. For further explanations we refer to chapters 22.3 and 22.4 of this book:

And this concludes the discussion of our first example. It has been kept quite informal. This is admissible because the goal was to explain an existing program. However, programming is designing, creating new programs. For this purpose, only a precise, formal description of our tool is adequate. In the next chapter, we introduce a formalism for the precise description of correct, "legal" program texts. This formalism makes it possible to determine in a rigorous manner whether a written text meets the language's rules.

3. A Notation to Describe the Syntax of Oberon

A formal language is an infinite set of sequences of symbols. The members of this set are called sentences, and in the case of a programming language these sentences are programs. The symbols are taken from a finite set called the vocabulary. Since the set of programs is infinite, it cannot be enumerated, but is instead defined by rules for their composition. Sequences of symbols that are composed according to these rules are said to be syntactically correct programs; the set of rules is the syntax of the language.

Programs in a formal language then correspond to grammatically correct sentences of spoken languages. Every sentence has a structure and consists of distinct parts, such as subject, object, and predicate. Similarly, a program consists of parts, called syntactic entities, such as statements, expressions, or declarations. If a construct A consists of B followed by C, i.e. the concatenation BC, then we call B and C syntactic factors and describe A by the syntactic formula

A = BC.

If, on the other hand, an A consists of a B or, alternatively, of a C, we call B and C syntactic terms and express A as

A = B | C.

Parentheses may be used to group terms and factors. It is noteworthy that here A, B, and C denote syntactic entities of the formal language to be described, whereas the symbols =, | , parentheses, and the period are symbols of the meta-notation describing syntax. The latter are called meta-symbols, and the meta-notation introduced here is called Extended Backus-Naur Formalism (EBNF).

In addition to concatenation and choice, EBNF also allows to express option and repetition. If a construct A may be either a B or nothing (empty), this is expressed as

A = [B].

and if an A consists of the concatenation of any number of Bs (including none), this is denoted by

A = {B}.

This is all there is to EBNF! A few examples show how sets of sentences are defined by EBNF formulas:

(A|B)(C|D)	AC AD BC BD
A[B]C	ABC AC
A{BA}	A ABA ABABA ABABABA ...
{A|B}C	C AC BC AAC ABC BBC BAC ...

Evidently, EBNF is itself a formal language. If it suits its purpose, it must at least be able to describe itself! In the following definition of EBNF in EBNF, we use the following names for entities:

statement:	a syntactic equation
expression:	a list of alternative terms
term:	a concatenation of factors
factor:	a single syntactic entity or a parenthesized expression

The formal definition of EBNF is now given as follows:

syntax =	{statement}.
statement =	identifier "=" expression ".".
expression =	term {"|" term}.
term =	factor {factor}.
factor =	identifier | string | "(" expression ")" | "[" expression "]" | "{" expression "}".

Identifiers denote syntactic entities; strings are sequences of symbols taken from the defined language's vocabulary. For the denotation of identifiers we adopt the widely used conventions for programming languages, namely:

An identifier consists of a sequence of letters and digits, where the first character must be a letter. A string consists of any sequence of characters enclosed by quote marks (or apostrophes).

A formal statement of these rules in terms of EBNF is given in the subsequent chapter.

4. Representation of Oberon Programs

The preceding chapter has introduced a formalism, by which the structures of well-formed programs will subsequently be defined. It defines, however, merely the way in which programs are composed as sequences of symbols, in contrast to sequences of characters. This "shortcoming" is quite intentional: the representation of symbols (and thereby programs) in terms of characters is considered too much dependent on individual implementations for the general level of abstraction appropriate for a language definition. The creation of an intermediate level of representation by symbol sequences provides a useful decoupling between language and ultimate program representation. The latter depends on the available character set. As a consequence, we need to postulate a set of rules governing the representation of symbols as character sequences. The symbols of the Oberon vocabulary are divided into the following classes:

identifiers, numbers, strings, operators and delimiters, and comments.

The rules governing their representation in terms of the standard ISO character set are the following:

1. *Identifiers* are sequences of letters and digits. The first character must be a letter. Capital and lower-case letters are considered as distinct.

identifier = letter {letter|digit}.

Examples of well-formed identifiers are

Alice   likely   jump   BlackBird   SR71

Examples of words which are no identifiers are

sound proof	(blank space is not allowed)
sound-proof	(neither is a hyphen)
2N	(first character must be a letter)
Miller's	(no apostrophe allowed)

Sometimes an identifier has to be qualified by another identifier; this is expressed by prefixing i with j and a period (j.i); the combined identifier is called a *qualified identifier* (abbreviated as *qualident*). Its syntax is

    qualident = {identifier "."} identifier.

2. *Numbers* are either integers or real numbers. The former are denoted by sequences of digits. Numbers must not include any spaces. Real numbers contain a decimal point and a fractional part. In addition, a scale factor may be appended. It is specified by the letter E and an integer which is possibly preceded by a sign. The E is pronounced as "times 10 to the power of".

Examples of well-formed numbers are

    1981   1   3.25   5.1E3   4.0E-10

Examples of character sequences that are not recognized as numbers are

    1,5   no comma may appear
    1'000'000   neither may apostrophs
    3.5En   no letters allowed (except the E)

The exact rules for forming numbers are given by the following syntax:

    number = integer | real.
    integer = digit {digit}.
    real = digit {digit} "." {digit} [ScaleFactor].
    ScaleFactor = "E" ["+"|"-"] digit {digit}.

Note: Integers are taken as octal numbers, if followed by the letter B, or as hexadecimal numbers if followed by the letter H.

3. *Strings* are sequences of any characters enclosed in quote marks. In order that the closing quote is recognized unambiguously, the string itself evidently cannot contain a quote mark.

Examples of strings are

    "no comment"
    "Buck's Corner"

4. *Operators and delimiters* are either special characters or reserved words. These latter are written in capital letters and must not be used as identifiers. Hence is it advantageous to memorize this short list of words. The operators and delimiters composed of special characters are

    +   addition, set union
    -           subtraction, set difference
    *   multiplication, set intersection
    /   division, symmetric set difference
    :=   assignment
    &   logical AND
    ~   logical NOT
    =   equal
    #   unequal
    <   less than
    >   greater than
    <=   less than or equal
    >=   greater than or equal
    ( )   parentheses
    [ ]   index brackets
    { }   set braces
    (* *)   comment brackets
    ^   dereferencing operator
    , . ; : .. | punctuation symbols

The reserved words are enumerated in the following list; their meaning will be explained throughout the subsequent chapters:

ARRAY	BEGIN	BY	CASE
CONST	DEFINITION	DIV	DO

ELSE	ELSIF	END	FOR
IF	IMPORT	IN	MOD
MODULE	OF	OR	POINTER
PROCEDURE	RECORD	REPEAT	RETURN
SET	THEN	TO	TYPE
UNTIL	VAR	WHILE	

It is customary to separate consecutive symbols by a space , i.e. one or several blanks. However, this is mandatory only in those cases where the lack of such a space would merge the two symbols into one. For example, in "IF x = y THEN" spaces are necessary in front of x and after y, but could be omitted around the equal sign.

5. Comments may be inserted between any two symbols. They are arbitrary sequences of characters enclosed in the comment brackets (* and *). Comments are skipped by compilers and serve as additional information to the human reader. They may also serve to signal instructions (options) to the compiler.

5. Statements and Expressions

The specification of an action is called a *statement*. Statements can be interpreted (executed), and that interpretation (execution) has an effect. The effect is a transformation of the state of the computation, the state being represented by the collective values of the program's variables. The most elementary action is the assignment of a value to a variable. The assignment statement has the form

    assignment = designator ":=" expression.

and its corresponding action consists of three parts in this sequence:

1. Evaluate the designator designating a variable.
2. Evaluate the expression yielding a value.
3. Replace the value of the variable identified in 1. by the value obtained in 2.

Simple examples of assignments are

    i := 1   x := y+z

Here i obtains the value 1, and x the sum of y and z, and the previous values are lost. Evidently, every variable in an expression must previously have been assigned a value. Observe that the following pairs of statements, executed in sequence, do not have the same effect:

    i := i+1; j := 2*i
    j := 2*i; i := i+1

Assuming the initial value i = 0, the first pair yields i = 1, j = 2, whereas the second pair yields j = 0. If we wish to exchange the values of variables i and j, the statement sequence

    i := j; j := i

will not have the desired effect. We must introduce a temporary value holder, say k, and specify the three consecutive assignments

    k := i; i := j; j := k

An *expression* is in general composed of several operands and operators. Its evaluation consists of applying the operators to the operands in a prescribed sequence, in general taking the operators from left to right. The operands may be constants, or variables, or functions. (The latter will be explained in a later chapter.) The identification of a variable will, in general, again require the evaluation of a designator. Here we shall confine our presentation to simple variables designated by an identifier. Arithmetic expressions (there exist other expressions too) involve numbers, numeric variables, and arithmetic operators. These include the basic operations of addition (+), subtraction (-), multiplication (*), and division. They will be discussed in detail in the chapter on basic data types. Here it may suffice to mention that the slash (/) is reserved for dividing real numbers, and that for integers we use the operator DIV which truncates the quotient.

An expression consists of consecutive terms. The expression

    T0 + T1 + ... + Tn

is equivalent to

((T0 + T1) + ...) + Tn

and is syntactically defined by the rules

SimpleExpression = ["+"|"-"] term {AddOperator term}.
AddOperator =   "+" | "-" | "OR".

Note: for the time being, the reader may consider the syntactic entities *expression* and *SimpleExpression* as equivalent. Their difference and the operators OR, AND, and NOT will be explained in the chapter on the data type BOOLEAN.

Each term similarly consists of factors. The term

F0 * F1 * ... * Fn

is equivalent to

((F0 * F1) * ...) * Fn

and is syntactically defined by the rules

term = factor {MulOperator factor}.
MulOperator = "*" | "/" | "DIV" | "MOD" | "&".

Each factor is either a constant, a variable, a function, or an expression itself enclosed by parentheses.

Examples of (arithmetic) expressions are

2 * 3 + 4 * 5	= (2*3)+(4*5)	= 26
15 DIV 4 * 4	= (15 DIV 4)*4	= 12
15 DIV (4*4)	= 15 DIV 16	= 0
2 + 3 * 4 - 5	= 2+(3*4)-5	= 9
6.25 / 1.25 + 1.5	= 5.0 + 1.5	= 6.5

The syntax of factors, implying that a factor may itself be an expression, is evidently recursive. The general form of designators will be explained later; here it suffices to know that an identifier denoting a variable or a constant is a designator.

factor = number | string | set | designator [ActualParameters] | "(" expression ")" | "~" factor.

The rules governing expressions are actually quite simple, and complicated expressions are rarely used. Nevertheless, we must point out a few basic rules that are well worth remembering.

1. Every variable in an expression must previously have been assigned a value.
2. Two operators must never be written side by side. For instance a * -b is illegal and must be written as a*(-b).
3. The multiplication sign must never be omitted when a multiplication is required. For example, 2n is illegal and must be written as 2*n.
4. MulOperators are binding more strongly than AddOperators.
5. When in doubt about evaluation rules (i.e. precedence of operators), use additional parentheses to clarify. For example, a + b * c may just as well be written as a+(b*c).

The assignment is but one of the possible forms of statements. Other forms will be introduced in the following chapters. We enumerate these forms by the following syntactic definition

statement = [assignment | ProcedureCall |
      WhileStatement | RepeatStatement | ForStatement |
      IfStatement | ReturnStatement].

Several of these forms are structured statements, i.e. some of their components may be statements again. Hence, the definition of statement is, like that of expressions, recursive.

The most basic structure is the sequence. A computation is a sequence of actions, where each action is specified by a statement, and is executed after the preceding action is completed. This strict sequentiality in time is an essential assumption of sequential programming. If a statement S1 follows S0, then we indicate this sequentiality by a semicolon

S0; S1

This statement separator (not terminator) indicates that the action specified by S0 is to be followed immediately by the action corresponding to S1. A sequence of statements is syntactically defined as

StatementSequence = statement {";" statement}.

The syntax of statements implies that a statement may consist of no symbols at all. In this case, the statement is said to be empty and evidently denotes the null action. This curiosity among statements has a definite reason: it allows semicolons to be inserted at places where they are actually superfluous, such as at the end of a statement sequence.

6. Control Structures

It is a prime characteristic of computers that individual actions can be selected, repeated, or performed conditionally depending on some previously computed results. Hence the sequence of actions performed is not always identical with the sequence of their corresponding statements. The sequence of actions is determined by control structures indicating repetition, selection, or conditional execution of given statements.

6.1 Repetitive Statements

The most common situation is the repetition of a statement or statement sequence under control of a condition: the repetition continues as long as the condition is met. This is expressed by the while statement. Its syntax is

WhileStatement = "WHILE" expression "DO" StatementSequence "END".

and its corresponding action is

1. Evaluate the condition which takes the form of an expression yielding the value TRUE or FALSE,
2. If the value is TRUE, execute the statement sequence and then repeat with step 1; if the value is FALSE, terminate.

The expression is of type BOOLEAN. This will be further discussed in the chapter on data types. Here it suffices to know that a simple comparison is a Boolean expression. An example was given in the introductory example, where repetition terminates when the two comparands have become equal. Further examples involving while statements are:

1. Initially, let $q = 0$ and $r = x$; then count the number of times y can be subtracted from x, i.e. compute the quotient $q = x$ DIV y, and remainder $r = x$ MOD y, if x and y are natural numbers.

WHILE r >= y DO r := r-y; q := q+1 END

2. Initially, let $z = 1$ and $i = k$; then multiply z k times by x, i.e. compute $z = x^k$, if z and k are natural numbers:

WHILE i > 0 DO z := z*x; i := i-1 END

When dealing with repetitions, it is important to remember the following points:

1. During each repetition, progress must be made towards meeting the goal, namely "getting closer" to satisfying the termination condition. An obvious corollary is that the condition must be somehow affected from within the repeated computation. The following statements are either incorrect or dependent on some critical precondition as stated.

```
WHILE i > 0 DO
  k := 2*k        (*i is not changed*)
END

WHILE i # 0 DO
  i := i-2        (*i must be even and positive*)
END

WHILE n # i DO
  n := n*i; i := i+1
END
```

2. If the condition is not satisfied initially, the statement is vacuous, i.e. no action is performed.

3. In order to obtain a grasp of the effect of the repetition, we need to establish a relationship that is stable, called an invariant. In the division example above, this is the equation $q*y + r = x$ holding each time the repetition is started. In the exponentiation example it is $z * x^i = x^k$ which, together with the termination condition $i = 0$ yields the desired result $z = x^k$.

4. The repetition of identical computations should be avoided (although a computer has infinite patience and will not complain). A simple rule is to avoid expressions within repetitive statements, in which no variable changes its value. For example, the statement

        WHILE i < 3*N DO tab[i] := x + y*z + z*i; i := i+1 END

should be formulated more effectively as

        n := 3*N; u := x + y*z;
        WHILE i < n DO tab[i] := u + z*i; i := i+1 END

In addition to the while statement, there is the repeat statement, syntactically defined as

        RepeatStatement = "REPEAT" StatementSequence "UNTIL" expression.

The essential difference to the while statement is that the termination condition is checked each time after (instead of before) execution of the statement sequence. As a result, the sequence is executed at least once. An advantage is that the condition may involve variables that are undefined when the repetition is started.

        REPEAT i := i+5; j := j+7; k := i DIV j UNTIL k > 23

        REPEAT r := r-y; q := q+1 UNTIL r < y

6.2 Conditional Statements

The conditional statement, also called if statement, has the form

        IfStatement =   "IF" expression "THEN" StatementSequence
                        {"ELSIF" expression "THEN" StatementSequence}
                        ["ELSE" StatementSequence]
                        "END".

The following example illustrates its general form.

        IF R1 THEN S1
        ELSIF R2 THEN S2
        ELSIF R3 THEN S3
        ELSE S4
        END

The meaning is obvious from the wording. However, it must be remembered that the expressions R1 ... R3 are evaluated one after the other, and that as soon as one yields the value TRUE, its corresponding statement sequence is executed, whereafter the IF statement is considered as completed. No further conditions are tested. Examples are:

        IF x = 0 THEN s := 0
        ELSIF x < 0 THEN s := -1
        ELSE s := 1
        END

        IF ODD(k) THEN z := z*x END

        IF k > 10 THEN k := k-10; d := 1 ELSE d := 0 END

The constructs discussed so far enable us to develop a few simple, complete programs as subsequently described. The first example is an extension of our introductory example computing the greatest common divisor *gcd* of two natural numbers x and y. The extension consists in the addition of two variables u and v, and statements which lead to the computation of the least common multiple (lcm) of x and y. The *lcm* and the *gcd* are related by the equation

        lcm(x,y) * gcd(x,y) = x*y

As explained earlier, this and following examples will be formulated as command procedures and are assumed to be embedded in module *Pattern* defined in Chap. 2 which imports modules *Texts* and *Oberon*.

```
PROCEDURE gcdlcm*;
  VAR x, y, u, v: INTEGER; S: Texts.Scanner;
BEGIN Texts.OpenScanner(S, Oberon.Par.text, Oberon.Par.pos);
  Texts.Scan(S); x := S.i; Texts.WriteString(W, "  x ="); Texts.WriteInt(W, x, 6);
  Texts.Scan(S); y := S.i; Texts.WriteString(W, "  y ="); Texts.WriteInt(W, y, 6);
  u := x; v := y;
  WHILE x # y DO
    (*gcd(x,y) = gcd(x0,y0), x*v + y*u = 2*x0*y0*)
    IF x > y THEN
      x := x - y; u := u + v
    ELSE
      y := y - x; v := v + u
    END
  END ;
  Texts.WriteInt(W, x, 6); Texts.WriteInt(W, (u+v) DIV 2, 6); Texts.WriteLn(W);
  Texts.Append(Oberon.Log, W.buf)
END gcdlcm.
```

This example again shows the nesting of control structures. The repetition expressed by a while statement includes a conditional structure expressed by an if statement, which in turn includes two statement sequences, each consisting of two assignments. This hierarchical structure is made transparent by appropriate indentation of the "inner" parts.

Another example demonstrating a hierarchical structure computes the real number x raised to the power i, where i is a non-negative integer.

```
PROCEDURE Power*;
  VAR i: INTEGER; x, z: REAL; S: Texts.Scanner;
BEGIN Texts.OpenScanner(S, Oberon.Par.text, Oberon.Par.pos); Texts.Scan(S);
  WHILE S.class = Texts.Real DO
    x := S.x; Texts.WriteString(W, "  x ="); Texts.WriteReal(W, x, 16);
    Texts.Scan(S); i := S.i; Texts.WriteString(W, "  i ="); Texts.WriteInt(W, i, 4);
    z := 1.0;
    WHILE i > 0 DO
      (* z * x^i = x0^i0 *)
      z := z*x; i := i-1
    END ;
    Texts.WriteReal(W, z, 16); Texts.WriteLn(W); Texts.Scan(S)
  END ;
  Texts.Append(Oberon.Log, W.buf)
END Power.
```

Here we subject the computation to yet another repetition: Each time a result has been computed, another value pair x, i is requested. This outermost repetition is controlled by the relation *S.class = Texts.Real*, which indicates whether a real number x had actually been read. As an example, activation of the command

        M.Power 5.0 2   2.0 5   3.0 3 ~

will generate the results 25.0, 32.0, 27.0.

The straight-forward computation of a power by repeated multiplication is, although obviously correct, not the most economical. We now present a more sophisticated and more efficient solution. It is based on the following consideration: The goal of the repetition is to reach the value i = 0. This is done by successively reducing i, while maintaining the invariant $z * x^i = x0^{i0}$, where x0 and i0 denote the initial values of x and i. A faster algorithm therefore must rely on decreasing i in larger steps. The solution given here halves i. But this is only possible, if i is even. Hence, if i is odd, it is first decremented by 1. Of course, each change of i must be accompanied by a corrective action on z in order to maintain the invariant. A detail: the subtraction of 1 from i is not expressed by an explicit statement, because it is performed implicitly by the subsequent division by 2. Two further details are noteworthy: The function

ODD(i) is TRUE, if i is an odd number, FALSE otherwise. x and z denote real values, as opposed to integer values. Hence, they can represent fractions, too.

```
PROCEDURE Power*;
  VAR i: INTEGER; x, z: REAL; S: Texts.Scanner;
BEGIN Texts.OpenScanner(S, Oberon.Par.text, Oberon.Par.pos); Texts.Scan(S);
  WHILE S.class = Texts.Real DO
    x := S.x; Texts.WriteString(W, " x ="); Texts.WriteReal(W, x, 16);
    Texts.Scan(S); i := S.i; Texts.WriteString(W, " i ="); Texts.WriteInt(W, i, 4);
    z := 1.0;
    WHILE i > 0 DO
      (* z * x^i = x0^i0 *)
      IF ODD(i) THEN z := z*x END ;
      x := x*x; i := i DIV 2
    END ;
    Texts.WriteReal(W, z, 16); Texts.WriteLn(W); Texts.Scan(S)
  END ;
  Texts.Append(Oberon.Log, W.buf)
END Power.
```

The next sample command has a structure that is almost identical to the preceding program. It computes the logarithm of a real number x whose value lies between 1 and 2. The invariant in conjunction with the termination condition (b = 0) implies the desired result sum = $\log_2(x)$.

```
PROCEDURE Log2*;
  VAR x, a, b, sum: REAL; S: Texts.Scanner;
BEGIN Texts.OpenScanner(S, Oberon.Par.text, Oberon.Par.pos); Texts.Scan(S);
  WHILE S.class = Texts.Real DO
    x := S.x; Texts.WriteReal(W, x, 15);   (*1.0 <= x < 2.0*)
    a := x; b := 1.0; sum := 0.0;
    REPEAT
      (*log2(x) = sum + b*log2(a)*)
      a := a*a; b := 0.5*b;
      IF a >= 2.0 THEN
        sum := sum + b; a := 0.5*a
      END
    UNTIL b < 1.0E-7;
    Texts.WriteReal(W, sum, 16); Texts.WriteLn(W); Texts.Scan(S)
  END ;
  Texts.Append(Oberon.Log, W.buf)
END Log2.
```

Normally, routines for computing standard mathematical functions need not be programmed in detail, because they are available from a collection of modules similar to those for input and output. Such a collection is, somewhat inappropriately, called a library. In the following example, again exhibiting the use of a repetitive statement, we use routines for computing the cosine and the exponential function from a library called *Math* and generate a table of values for a damped oscillation. Typically, the available standard routines include the sin, cos, exp, ln (logarithm), sqrt (square root), and the arctan functions.

```
PROCEDURE Oscillation*;
  CONST dx = 0.19634953;  (*pi/16*)
  VAR i, n: INTEGER;
    x, y, r: REAL;  S: Texts.Scanner;
BEGIN Texts.OpenScanner(S, Oberon.Par.text, Oberon.Par.pos);
  Texts.Scan(S); n := S.i; Texts.WriteInt(W, n, 6);
  Texts.Scan(S); r := S.x; Texts.WriteReal(W, r, 15); Texts.WriteLn(W);
  i := 0; x := 0.0;
  REPEAT x := x + dx; i := i+1;
    y := Math.exp(-r*x) * Math.cos(x);
    Texts.WriteReal(W, x, 15); Texts.WriteReal(W, y, 15); Texts.WriteLn(W)
  UNTIL i >= n;
  Texts.Append(Oberon.Log, W.buf)
```

END Oscillation.

7. Elementary Data Types

We have previously stated that all variables need be declared. This means that their names are introduced in the heading of the program. In addition to introducing the name (and thereby enabling a compiler to detect and indicate misspelled identifiers), declarations have the purpose of associating a data type with each variable. This data type represents information about the variable which is permanent in contrast, for example, to its value. This information may again aid in detecting inconsistencies in an erroneous program, inconsistencies that are detectable by mere inspection of the program text without requiring its interpretation.

The type of a variable determines its set of possible values and the operations which may be applied to it. Each variable has a single type that may be deduced from its declaration. Each operator requires operands of a specific type and produces a result of a specific type. It is therefore visible from the program text, whether or not a given operator is applicable to a given variable.

Data types may be declared in the program. Such constructed types are usually based on composition of basic types. There exist a number of most frequently used, elementary types that are basic to the language and need not be declared. They are called standard types and will be introduced in this chapter, although some have already appeared in previous examples.

In fact, not only variables have a type, but so do constants, functions, operands and (results of) operators. In the case of a constant, the type is usually deducible from the constant's notation, otherwise from its explicit declaration.

First we describe the standard data types of Oberon, and thereafter elaborate on the form of declarations of variables and constants. Further kinds of data types and declarations are deferred to later chapters.

7.1 The Type INTEGER

This type represents the whole numbers, and any value of type INTEGER is therefore an integer. Operators applicable to integers include the basic arithmetic operations

+	addition
-	subtraction
*	multiplication
DIV	division
MOD	modulus (remainder)

Integer division is denoted by DIV. If we define the integer quotient and the modulus of x and y by

$$q = x \ DIV \ y, \quad r = x \ MOD \ y$$

then q and r are related by the equation $x = q*y + r$ and by the constraint $0 <= r < y$. For example

15 DIV 4 = 3	15 MOD 4 = 3	15 = 3*4 + 3
-15 DIV 4 = -4	-15 MOD 4 = 1	-15 = (-4)*4 + 1

Sign inversion is denoted by the monadic minus sign. Furthermore, there exist the operators ABS(x) and ODD(x), the former yielding the absolute value of x, the latter the Boolean result "x is odd".

Every computer will restrict the set of values of type INTEGER to a finite set of integers, usually the interval $-2^{N-1} \ ... \ 2^{N-1}-1$, where N is a small integer, often 16 or 32, depending on the number of bits a computer uses to represent an integer. If an arithmetic operation produces a result that lies outside that interval, then overflow is said to have occured. The computer will give an appropriate indication and, usually, terminate the computation. The programmer should ensure that overflows will not result during execution of the program.

7.2 The Type REAL

Values of type REAL are real numbers. The available operators are again the basic arithmetic operations and ABS. Division is denoted by / (instead of DIV). Constants of type REAL are characterized by having a decimal point and possibly a decimal scale factor. Examples of real number denotations are

1.5   1.50   1.5E2   2.34E-2   0.0

The scale factor consists of the capital letter E followed by an integer. It means that the preceding real number is to be multiplied by 10 raised to the scale factor. Hence

1.5E2 = 150.0,   2.34E-2 = 0.0234

The important point to remember is that real values are internally represented as pairs consisting of a fractional number and a scale factor. This is called *floating-point* representation. Of course, both parts consist of a finite number of digits. As a consequence, the representation is inherently inexact, and computations involving real values are inexact because each operation may be subject to truncation or rounding.

The following program makes that inherent imprecision of computations involving operands of type REAL apparent. It computes the harmonic function

$H(n) = 1 + 1/2 + 1/3 + ... + 1/n$

in two different ways, namely once by summing terms from left to right, once from right to left. According to the rules of arithmetic, the two sums ought to be equal. However, if values are truncated (or even rounded), the sums will differ for sufficiently large n. The correct way is evidently to start with the small terms.

```
PROCEDURE Harmonic*;
  VAR i, n: INTEGER;
      x, d, s1, s2: REAL; S: Texts.Scanner;
BEGIN Texts.OpenScanner(S, Oberon.Par.text, Oberon.Par.pos); Texts.Scan(S);
  WHILE S.class = Texts.Int DO
    n := S.i; Texts.WriteString(W, "n = "); Texts.WriteInt(W, n, 6);
    s1 := 0.0; d := 0.0; i := 0;
    REPEAT
       d := d + 1.0; INC(i);
       s1 := s1 + 1.0/d;
    UNTIL i >= n;
    Texts.WriteReal(W, s1, 16);
    s2 := 0.0;
    REPEAT
       s2 := s2 + 1.0/d;
       d := d - 1.0; DEC(i);
    UNTIL i = 0;
    Texts.WriteReal(W, s2, 16); Texts.WriteLn(W); Texts.Scan(S)
  END ;
  Texts.Append(Oberon.Log, W.buf)
END Harmonic.
```

The major reason for strictly distinguishing between real numbers and integers lies in the different representation used internally. Hence, also the arithmetic operations are implemented by instructions which are distinct for each type. Oberon postulates that expressions with mixed operands always are of type REAL.

A real number x is transformed into an integer by the standard function ENTIER(x). It yields the largest integer not greater than x.

ENTIER(3.1416) = 3   ENTIER(-2.7) = -3   ENTIER(5.0) = 5

7.3 The Type BOOLEAN

A BOOLEAN value is one of the two logical truth values denoted by the standard identifiers TRUE and FALSE. Boolean variables are usually denoted by identifiers which are adjectives, the value TRUE implying the presence, FALSE the absence of the indicated property. A set of logical operators is provided which, together with BOOLEAN variables, form BOOLEAN expressions. These operators are & (and), OR, and ~ (not). Their results are explained as follows

```
p & q      = "both p and q are TRUE"
p OR q     = "either p or q or both are TRUE"
~p         = "p is FALSE"
```

The operators' exact definition, however, is slightly different, although the results are identical:

    p & q   = IF p THEN q ELSE FALSE
    p OR q   = IF p THEN TRUE ELSE q

This definition implies that the second operand need not be evaluated, if the result is already known from the evaluation of the first operand. The notable property of this definition of Boolean connectives is that their result may be well-defined even in cases where the second operand is undefined. As a consequence, the order of the operands may be significant.

Sometimes it is possible to simplify Boolean expressions by application of simple transformation rules. A particularly useful rule is de Morgan's law stating the equivalences

    ~p & ~q = ~(p OR q)
    ~p OR ~q = ~(p & q)

Relations produce a result of type BOOLEAN, i.e. TRUE if the relation is satisfied, FALSE if not. For example

    7 = 12   FALSE
    7 < 12   TRUE
    1.5 >= 1.6   FALSE

Relations are syntactically classified as expressions, and the two comparands are so-called SimpleExpressions (see also chapter on expressions and statements). The result of a comparison is of type BOOLEAN and can be used in control structures such as if, while, and repeat statements. The symbol # stands for unequal.

    expression = SimpleExpression [relation SimpleExpression].
    relation   = "=" | "#" | "<" | "<=" | ">" | ">=" | "IN".

It should be noted that, similar to arithmetic operators, there exists a precedence hierarchy among the Boolean operators. ~ has the highest precedence, then follows &, then OR, and last are the relations. As with arithmetic expressions, parentheses may be used freely to make the association of operators explicit. Examples of Boolean expressions are

    x = y
    (x <= y) & (y < z)
    (x > y) OR (y >= z)
    ~p OR q

Note that a construct such as x < y & z < w is illegal.

The example above draws attention to the rule that the operands of an operator (including relational operators) must be of the same type. The following relations are therefore illegal:

    1   = TRUE
    5   = 5.0
    i+j = p OR q

Also incorrect is e.g. x <= y < z, which must be expanded into (x <= y) & (y < z). However, the following are correct Boolean expressions:

    i+j < k-m
    p OR q = (i<j)

A final hint: although "p = TRUE" is legal, it is considered poor style and better expressed as p. Similarly, replace "p = FALSE" by "~p".

7.4 The Type CHAR

Every computer system communicates with its environment via some input and output devices. They read, write, or print elements taken from a fixed set of characters. This set constitutes the value range of the type CHAR. Unfortunately, different brands of computers may use different character sets, which makes communication between them (i.e. the exchange of programs and data) difficult and often tedious. However, there exist internationally standardised sets. Here we refer to ISO's Latin-1 set.

The Latin-1 standard defines a set of 128 characters, 33 of which are so-called (non-printing) *control characters*. The remaining 95 elements are visible printing characters shown in the following table.

The set is ordered, and each character has a fixed position or ordinal number. For example, A is the 66th character and has the ordinal number 65. The ISO standard, however, leaves a few places open, and they can be filled with different characters according to national desires to establish national standards. Most widely used is the American standard, also called ASCII (American Standard Code for Information Interchange). Here we tabulate the ASCII set. The ordinal number of a character is obtained by taking its row number and adding its column number. These numbers are customarily given in hexadecimal form, and we follow this habit here too. The first two columns contain the control characters; they are customarily denoted by abbreviations hinting at their intended meaning. However, this meaning is not inherent in the character code, but only defined by its interpretation. At this point it suffices to remember that these characters are (usually) not printable.

Table of ASCII characters

	0	10	20	30	40	50	60	70
0	nul	dle		0	@	P	`	p
1	soh	dc1	!	1	A	Q	a	q
2	stx	dc2	"	2	B	R	b	r
3	etx	dc3	#	3	C	S	c	s
4	eot	dc4	$	4	D	T	d	t
5	enq	nak	%	5	E	U	e	u
6	ack	syn	&	6	F	V	f	v
7	bel	etb	'	7	G	W	g	w
8	bs	can	(	8	H	X	h	x
9	ht	em	)	9	I	Y	i	y
A	lf	sub	*	:	J	Z	j	z
B	vt	esc	+	;	K	[	k	{
C	ff	fs	,	<	L	\	l	|
D	cr	gs	-	=	M	]	m	}
E	so	rs	.	>	N	^	n	~
F	si	us	/	?	O	_	o	del

Constants of type CHAR are denoted by the character being enclosed in quote marks (single-character string). They can be assigned to variables of type CHAR. These values cannot be used in arithmetic operations. Arithmetic operators can be applied, however, to their ordinal numbers obtained from the transfer function ORD(ch). Inversely the character with the ordinal number n is obtained with the transfer function CHR(n). These two complementary functions are related by the equations

$$CHR(ORD(ch)) = ch \quad and \quad ORD(CHR(n)) = n$$

for $0 <= n < 128$. They permit to compute the numeric value represented by the digit ch as

$$ORD(ch) - ORD("0")$$

and to compute the digit representing the numeric value n as

$$CHR(n + ORD("0"))$$

These two formulas depend on the adjacency of the 10 digits in the ISO character set, where ORD("0") = 48. They are typically used in routines converting sequences of digits into numbers and vice-versa. The following program piece reads digits and assigns the decimally represented number to a variable x.

```
x := 0; Texts.Read(R, ch);
WHILE ("0" <= ch) & (ch <= "9") DO
    x := 10*x + (ORD(ch) - ORD("0")); Texts.Read(R, ch)
END
```

Control characters are used for various purposes, mainly to control the functions of devices, but also to delineate and structure text. An important role of control characters is to specify the end of a line or of a page of text. There is no universally accepted standard for this purpose. We shall denote the control character signifying a line end by the identifier EOL (end of line); its actual value depends on the system used.

In order to be able to denote non-printable characters, Oberon uses their hexadecimal ordinal number followed by the capital letter X. For example, 0DX is the value of type CHAR denoting the control character *cr* (carriage return, line end) with ordinal number 13.

7.5 The Type SET

The values which belong to the type SET are sets of integers between 0 and N-1, where N is a constant defined by the computer system used. It is usually the computer's wordlength or a small multiple of it, typically 32. Constants of this type are denoted as sets. Examples are

$\{2, 3, 5, 7, 11\}$ $\{0\}$ $\{8 .. 15\}$ $\{\}$

The notation m .. n is a shorthand for m, m+1, ... , n-1, n.

```
set = [qualident] "{" [element {"," element}] "}".
element = expression [".." expression].
```

Operations on sets are:

+   set union
-   set difference
*   set intersection
/   symmetric set difference

Assuming that i denotes an element and u, v denote sets, these operations are defined in terms of set membership as

1. i IN (u+v) = (i IN u) OR (i IN v), i.e. the integer i is in the union u+v, if either in u or in v (or in both).
2. i IN (u-v) = (i IN u) AND NOT (i IN v), i.e. i is in the difference u-v, if it is in u, but not in v.
3. i IN (u*v) = (i In u) AND (i IN v), i.e. i is in the intersection u*v, if it is in both u and v.
4. i IN (u/v) = ((i IN u) # (i IN v)), i.e. i is in the symmetric difference u/v, if it is in either u or in v, but not in both.

The membership operator IN is regarded as a relational operator. The expression i IN u is of type BOOLEAN. It is TRUE, if i is a member of the set u. Sets are represented in computer systems as sets of bits, i.e. by the characteristic function of the set. The i'th bit of u, for example, is 1, if i is a member of u, 0 otherwise. Hence, the set operators are implemented as logical operations applied to the N members of the set variable. They are therefore very efficient and their execution time is usually even less than that of an addition of integers.

8. Constant and Variable Declarations

It has been mentioned previously that all identifiers used in a program must be declared in the program's heading, unless they are standard identifiers known in every program (or are imported from some other module).

If an identifier is to denote a constant value, it must be introduced by a constant declaration which indicates the value for which the constant identifier stands. A constant declaration has the form

```
ConstantDeclaration = identifier "=" ConstExpression.
ConstExpression = expression
```

A ConstExpression is an expression containing constants only. More precisely, its evaluation must be possible by a mere textual scan without execution of the program A sequence of constant declarations is preceded by the symbol CONST. Example:

```
CONST N = 16;
        EOL = 0DX;
        empty = {};
        M = N-1;
```

Constants with explicit names aid in making a program readable, provided that the constants are given suggestive names. If, e.g., the identifier N is used instead of its value throughout a program, a change of that constant can be achieved by changing the program in a single place only, namely in the declaration of N. This avoids the common mistake that some instances of the constant, spread over the entire program text, remain undetected and therefore are not updated, leading to inconsistencies.

A variable declaration looks similar to a constant declaration. The place of the constant's value is taken by the variable's type which, in a sense, can be regarded as the variable's constant property. Instead of an equal sign, a colon is used.

```
VariableDeclaration = IdentList ":" type.
IdentList = identifier {"," identifier}.
```

Variables of the same type may be listed in the same declaration, and a sequence of declarations is preceded by the symbol VAR. Example:

```
VAR i, j, k: INTEGER;
  x, y, z: REAL;
  ch: CHAR
```

9. The Data Structure Array

So far, we have given each variable an individual name. This is impractical, if many variables are necessary that are treated in the same way and are of the same type, such as, for example, if a table of data is to be constructed. In this case, we wish to give the entire set of data a name and to denote individual elements with an identifying number, a so-called *index*. The data type is then said to be structured - more precisely: array structured. In the following example, the variable *a* consists of N elements, each being of type INTEGER, and the indices range from 0 to N-1.

```
VAR a: ARRAY N OF INTEGER
```

An element is then designated by the array's identifier followed by its selecting index, e.g. a[i], where i is an expression whose value must lie within the index range specified in the array's declaration. Syntactically, a[i] is a designator, and the expression i is the selector. If, for example, all N elements of *a* are to be given the value 0, this can be expressed conveniently by a repetitive statement, where the index is given a new value each time.

```
i := 0;
REPEAT a[i] := 0; INC(i) UNTIL i = N
```

INC(i) is synonymous with i := i+1. This example illustrates a situation that occurs so frequently that Oberon provides a special control structure which expresses it more concisely. It is called the *for statement*:

```
FOR i := 0 TO N-1 DO a[i] := 0 END
```

Its general form is

```
ForStatement =
  "FOR" identifier ":=" expression "TO" expression ["BY" ConstExpression] "DO"
    StatementSequence
  "END".
```

The expressions before and after the symbol TO define the range through which the so-called *control variable* (i) progresses. An optional parameter determines the incrementing (decrementing) value. If it is omitted, 1 is assumed as default value.

It is recommended that the for statement be used in simple cases only; in particular, no components of the expressions determining the range must be affected by the repeated statements, and, above all, the control variable itself must not be changed by the repeated statements. The value of the control variable must be considered as undefined after the for statement is terminated.

Some additional examples should demonstrate the use of array structures and the for statement. In the first one, the sum of the N elements of array *a* is to be computed:

```
sum := 0;
FOR i := 0 TO N-1 DO sum := a[i] + sum END
```

In the second example, the minimum value is to be found, and also its index. The repetition's invariant is min = minimum(a[0], ... , a[i-1]).

```
min := a[0]; k := 0;
FOR i := 1 TO N-1 DO
  IF a[i] < min THEN k := i; min := a[k] END
END
```

In the third example, we make use of the second to sort the array in ascending order:

```
FOR i := 0 TO N-2 DO
  min := a[i]; k := i;
  FOR j := i TO N-1 DO
    IF a[j] < min THEN k := j; min := a[k] END
  END;
  a[k] := a[i]; a[i] := min
END
```

The for statement is evidently appropriate, if all elements within a given index range are to be processed. It is inappropriate in most other cases. If, for example, we wish to find the index of an element equal to a given value x, we have no advance knowledge how many elements have to be inspected. Hence, the use of a while (or repeat) statement is recommended. This algorithm is known as linear search.

```
i := 0;
WHILE (i < N) & (a[i] # x) DO  INC(i) END
```

From the negation of the continuation condition, applying de Morgan's law, we infer that upon termination of the while statement the condition (i = N) OR (a[i] = x) holds. If the latter term is TRUE, the desired element is found and i is its index; if i = N, no a[i] equals x.

We draw attention to the fact that the termination condition is a composite one, and that it is possible to simplify this condition by a common technique. Remember, that repetition must terminate, either if the desired element is found, or if the array's end is reached. The "trick" now consists in marking the end by a sentinel, namely the value x itself, which will automatically stop the search. All that is required is the addition of a dummy element a[N] at the array's end, serving as sentinel:

```
a: ARRAY N+1 OF INTEGER;
```

```
a[N] := x; i := 0;
WHILE a[i] # x DO INC(i) END
```

If upon termination i = N, no original element has the value x, otherwise i is the desired index.

A more challenging problem is the search for a desired element with value x in an array that is ordered, i.e. a[i-1] <= a[i] for all i = 1 ... N-1. The best technique is the so-called *binary search*: inspect the middle element, then apply the same method to either the left or right half. This is expressed by the following piece of program, assuming N > 0. The repetition's invariant is

```
a[k] < x   for all k = 0 ... i-1    and    a[k] > x   for all k = j+1 ... N-1
```

```
i := 0; j := N-1; found := FALSE;
REPEAT mid := (i+j) DIV 2;
  IF x < a[mid] THEN j := mid -1
  ELSIF x > a[mid] THEN i := mid+1
  ELSE found := TRUE
UNTIL (i > j) OR found
```

Because each step halves the interval in which x is searched, the number of needed comparisons is only $\log_2 N$. A somewhat more efficient version which avoids the composite termination condition is

```
i := 0; j := N-1;
REPEAT mid := (i+j) DIV 2;
  IF x <= a[mid] THEN j := mid-1 END;
  IF x >= a[mid] THEN i := mid+1 END
UNTIL i > j;
IF i > j+1 THEN found ELSE not found END
```

An even more sophisticated version is given below. Its ingenious idea is not to terminate immediately when the element is found, which is rare compared to the number of unsuccessful comparisons.

```
i := 0; j := N;
REPEAT mid := (i+j) DIV 2;
  IF x < a[mid] THEN j := mid ELSE i := mid+1 END
UNTIL i >= j;
IF j < N) & (a[j] = x) THEN (*found*) … END
```

This ends our list of examples of typical uses of simple arrays.

The elements of an array are all of the same type, but this type may again be an array (in fact, it may be any structure, as will be seen later). An array of arrays is called a multidimensional array or *matrix*, because each index may be considered as spanning a dimension in a Cartesian space. Examples of two-dimensional arrays are

    a: ARRAY N, N OF REAL
    T: ARRAY M, N OF CHAR

These are actually abbreviations of the full forms

    a: ARRAY N OF ARRAY N OF REAL
    T: ARRAY M OF ARRAY N OF CHAR

The general syntax of an array type is

    ArrayType = "ARRAY" ConstExpression {"," ConstExpression } "OF" type.

where the expression denotes the length of the array.

The syntax of designators allows for a similar abbreviation as used in declarations, namely a[i, j] in place of a[i][j]. The latter form expresses more clearly that j is the selector on the array a[i]. The syntax for the array element designator is

    designator = qualident {"[" ExpList "]"}.
    ExpList = expression {"," expression}.

In uses of matrices, the for statement comes to its full bloom, particularly in numeric applications. The canonical example is the multiplication of two matrices, where each element of the product c = a*b is defined as

    c[i, j] = a[i, 0] * b[0, j] + a[i, 1] * b[1, j] + ... + a[i, K-1] * b[K-1, j]

Given the declarations

    a: ARRAY M, K OF REAL;
    b: ARRAY K, N OF REAL;
    c: ARRAY M, N OF REAL

the multiplication algorithm consists of three nested repetitions as follows:

```
FOR i := 0 TO M-1 DO
  FOR j := 01 TO N-1 DO
    sum := 0.0;
    FOR k := 0 TO K-1 DO
      sum := a[i, k]*b[k, j] + sum
    END;
    c[i,j] := sum
  END
END
```

In a second example we demonstrate the search of a word in a table, so-called *table search*, where each word is an array of characters. We assume the table is declared by the T given in the example above, and that x is given as

    x: ARRAY N OF CHAR

Our solution employs the typical linear search

```
i := 0; found := FALSE;
WHILE ~found & (i < M) DO
  found := T[i] = x;
  INC(i)
END
```

If we define equality between two words x and y as x[j] = y[j] for all j = 0 ... N-1, then the "inner" search can be expressed as

```
    j := 0; equal := TRUE;
    WHILE equal & (j < N) DO
       equal := T[i, j] = x[j]; INC(j)
    END;
    found := equal
```

This solution appears to be somewhat cumbersome. It is transformable into a simpler form, if M > 0 and N > 0. The complete table search is then expressible as

```
    i := 0;
    REPEAT j := 0;
       REPEAT b := T[i, j] # x[j]; INC(j)
       UNTIL b OR (j = N);
       INC(i)
    UNTIL ~b OR (i = M)
```

The result *b* means "the word x has not been found".

We have now laid enough ground work to develop meaningful, entire programs and shall present three examples, all of them involving arrays.

In the first example, the goal is to generate a list of powers of 2, each line showing the values 2^i, i, and 2^{-i}. This task is quite simple, if the type REAL is used. The program then contains the core

```
    d := 1; f := 1.0;
    FOR i := 1 TO N DO
       d := 2*d; write(d);   (* d = 2^i *)
       write(i);
       f := f/2.0; write(f)    (* f = 2^(-i) *)
    END
```

However, our task shall be to generate exact results with as many decimal digits as needed. For this reason, we present both the whole number $d = 2^i$ and the fraction $f = 2^{-i}$ by arrays of digits, each in the range 0 ... 9. For f we require N , for d only $\log_{10}N$ digits. Note that the doubling of *d* proceeds from right to left, the halving of *f* from left to right. The table of results is shown below.

```
    PROCEDURE PowersOf2*;
       CONST M = 11; N = 32; (*M ~ N*log(2) *)
       VAR i, j, k, exp, c, r, t: INTEGER;
          d: ARRAY M OF INTEGER;
          f: ARRAY N OF INTEGER;
    BEGIN d[0] := 1; k := 1;
       FOR exp := 1 TO N-1 DO
          (* compute d = 2^exp *)  c := 0; (*carry*)
          FOR i := 0 TO k-1 DO
             t := 2*d[i] + c;
             IF t >= 10 THEN d[i] := t - 10; c := 1 ELSE d[i] := t; c := 0 END
          END ;
          IF c > 0 THEN d[k] := 1; INC(k) END ;
          (*output d[k-1] ... d[0]*) i := M;
          REPEAT DEC(i); Texts.Write(W, " ") UNTIL i = k;
          REPEAT DEC(i); Texts.Write(W, CHR(d[i]+ORD("0"))) UNTIL i = 0;
          Texts.WriteInt(W, exp, 4);
          (*compute and output f = 2^(-exp) *)
          Texts.WriteString(W, "  0."); r := 0; (*remainder*)
          FOR j := 1 TO exp-1 DO
             r := 10*r + f[j]; f[j] := r DIV 2;
             r := r MOD 2; Texts.Write(W, CHR(f[j]+ORD("0")))
          END ;
          f[exp] := 5; Texts.Write(W, "5"); Texts.WriteLn(W)
       END ;
       Texts.Append(Oberon.Log, W.buf)
    END PowersOf2.
```

Output of program PowersOf2:

```
          2   1   0.5
          4   2   0.25
          8   3   0.125
         16   4   0.0625
         32   5   0.03125
         64   6   0.015625
        128   7   0.0078125
        256   8   0.00390625
        512   9   0.001953125
       1024  10   0.0009765625
       2048  11   0.00048828125
       4096  12   0.000244140625
       8192  13   0.0001220703125
      16384  14   0.00006103515625
      32768  15   0.000030517578125
      65536  16   0.0000152587890625
     131072  17   0.00000762939453125
     262144  18   0.000003814697265625
     524288  19   0.0000019073486328125
    1048576  20   0.00000095367431640625
    2097152  21   0.000000476837158203125
    4194304  22   0.0000002384185791015625
    8388608  23   0.00000011920928955078125
   16777216  24   0.000000059604644775390625
   33554432  25   0.0000000298023223876953125
   67108864  26   0.00000001490116119384765625
  134217728  27   0.000000007450580596923828125
  268435456  28   0.0000000037252902984619140625
  536870912  29   0.00000000186264514923095703125
 1073741824  30   0.000000000931322574615478515625
 2147483648  31   0.0000000004656612873077392578125
 4294967296  32   0.00000000023283064365386962890625
```

Our second example is similar in nature. Its task is to compute the fractions d = 1/i exactly. The difficulty lies, of course, in the representation of those fractions that are infinite sequences of digits, e.g. 1/3 = 0.333.... Fortunately, all fractions have a repeating period, and a reasonable and useful solution is to mark the beginning of the period and to terminate at its end. How do we find the beginning and the end of the period? Let us first consider the algorithm for computing the digits of the fraction.

Starting out with *rem* = 1, we repeat multiplying by 10 and dividing the product by i. The integer quotient is the next digit and the remainder is the new value of *rem*. This is precisely the conventional method of division, as illustrated by the following piece of program and the example with i = 7:

```
     1.000000 / 7 = 0.142857
    1 0
       30
        20
         60
          40
           50
            1
```

```
rem := 1;
REPEAT rem := 10 * rem; nextDigit := rem DIV i; rem := rem MOD i UNTIL ...
```

We know that the period has ended as soon as a remainder occurs which had been encountered previously. Hence, our recipe is to remember all remainders and their indices. The latter designate the place where the period had started. We denote these indices by x and give elements of x initial value 0. In the above explained division by 7, the values of x are

x[1] = 1, x[2] = 3, x[3] = 2, x[4] = 6, x[5] = 4, x[6] = 5

```
PROCEDURE Fractions*;
  CONST Base = 10; N = 32;
  VAR i, j, m, rem: INTEGER;
     d: ARRAY N OF INTEGER; (*digits*)
     x: ARRAY N+1 OF INTEGER; (*index*)
BEGIN
  FOR i := 2 TO N DO
```

```
    FOR j := 0 TO i-1 DO x[j] := 0 END ;
    m := 0; rem := 1;
    REPEAT m := m+1; x[rem] := m;
      rem := Base * rem; d[m] := rem DIV i; rem := rem MOD i
    UNTIL x[rem] # 0;
    Texts.WriteInt(W, i, 6); Texts.WriteString(W, "  0.");
    FOR j := 1 TO x[rem]-1 DO Texts.Write(W, CHR(d[j]+ORD("0"))) END ;
    Texts.Write(W, "'");
    FOR j := x[rem] TO m DO Texts.Write(W, CHR(d[j]+ORD("0"))) END ;
    Texts.WriteLn(W)
  END ;
  Texts.Append(Oberon.Log, W.buf)
END Fractions.
```

Output of program Fractions:

```
      2   0.5'0
      3   0.'3
      4   0.25'0
      5   0.2'0
      6   0.1'6
      7   0.'142857
      8   0.125'0
      9   0.'1
     10   0.1'0
     11   0.'09
     12   0.08'3
     13   0.'076923
     14   0.0'714285
     15   0.0'6
     16   0.0625'0
     17   0.'0588235294117647
     18   0.0'5
     19   0.'052631578947368421
     20   0.05'0
     21   0.'047619
     22   0.0'45
     23   0.'0434782608695652173913
     24   0.041'6
     25   0.04'0
     26   0.0'384615
     27   0.'037
     28   0.03'571428
     29   0.'0344827586206896551724137931
     30   0.0'3
     31   0.'032258064516129
     32   0.03125'0
```

Our last example of a program computes a list of prime numbers. It is based on the idea of inspecting the divisibility of successive integers. The tested integers are obtained by incrementing alternatively by 2 and 4, thereby avoiding multiples of 2 and 3 ab initio. Divisibility needs to be tested for prime divisors only, which are obtained by storing previously computed results.

```
    PROCEDURE Primes*;
      CONST N = 252; M = 16;  (*M ~ sqrt(N)*)
        LL = 10; (*no. of primes placed on a line*)
      VAR i, k, x: INTEGER;
        inc, lim, square, L: INTEGER;
        prime: BOOLEAN;
        P,V: ARRAY M+1 OF INTEGER;
    BEGIN L := 0; x := 1; inc := 4; lim := 1; square := 9;
      FOR i := 3 TO N DO
        (* find next prime number p[i] *)
        REPEAT x := x + inc; inc := 6 - inc;
          IF square <= x THEN
            INC(lim); V[lim] := square; square := P[lim+1] * P[lim+1]
          END ;
          k := 2; prime := TRUE;
          WHILE prime & (k < lim) DO
```

```
      INC(k);
      IF V[k] < x THEN V[k] := V[k] + 2*P[k] END ;
      prime := x # V[k]
    END
  UNTIL prime;
  IF i <= M THEN P[i] := x END ;
  Texts.WriteInt(W, x, 6); L := L+1;
  IF L = LL THEN Texts.WriteLn(W); L := 0 END
 END ;
 Texts.Append(Oberon.Log, W.buf)
END Primes.
```

Output of program Primes:

```
   5     7    11    13    17    19    23    29    31    37
  41    43    47    53    59    61    67    71    73    79
  83    89    97   101   103   107   109   113   127   131
 137   139   149   151   157   163   167   173   179   181
 191   193   197   199   211   223   227   229   233   239
 241   251   257   263   269   271   277   281   283   293
 307   311   313   317   331   337   347   349   353   359
 367   373   379   383   389   397   401   409   419   421
 431   433   439   443   449   457   461   463   467   479
 487   491   499   503   509   521   523   541   547   557
 563   569   571   577   587   593   599   601   607   613
 617   619   631   641   643   647   653   659   661   673
 677   683   691   701   709   719   727   733   739   743
 751   757   761   769   773   787   797   809   811   821
 823   827   829   839   853   857   859   863   877   881
 883   887   907   911   919   929   937   941   947   953
 967   971   977   983   991   997  1009  1013  1019  1021
1031  1033  1039  1049  1051  1061  1063  1069  1087  1091
1093  1097  1103  1109  1117  1123  1129  1151  1153  1163
1171  1181  1187  1193  1201  1213  1217  1223  1229  1231
1237  1249  1259  1277  1279  1283  1289  1291  1297  1301
1303  1307  1319  1321  1327  1361  1367  1373  1381  1399
1409  1423  1427  1429  1433  1439  1447  1451  1453  1459
1471  1481  1483  1487  1489  1493  1499  1511  1523  1531
1543  1549  1553  1559  1567  1571  1579  1583  1597  1601
```

These examples show that arrays are a fundamental feature used in most programs. There exist hardly any programs of relevance outside the classroom which do not employ repetitions and arrays (or analogous data structures).

Part 2

10. Procedures

Consider the task of processing a set of data consisting of a header and a sequence of N similar individual units. It might be generally described as

```
ReadHeader;
ProcessHeader;
WriteHeader;
FOR i := 1 TO N DO
   ReadUnit; ProcessUnit;
   Write(i); WriteUnit
END
```

Clearly, the description of the original task has been made in terms of subtasks, emphasizing the dominant structure and supressing details. Of course, the subtasks *ReadHeader, ProcessHeader*, etc. must now be further described with all the necessary details. Instead of replacing these descriptive English words with elaborate Oberon programs, we may consider these words as identifiers and define the details of the subtasks by textually separate pieces of program, called *procedures* (or subroutines). These definitions are called *procedure declarations*, because they define the actions of the procedure and give it a name. The identifiers in the main program referring to these declarations are said to be *procedure calls*, and their action is to invoke the procedure. Syntactically, the procedure call is a statement.

Procedures play a fundamental role in program design. They aid in displaying the algorithm's structure and in decomposing a program into logically coherent units. This is particularly important in the case of complex algorithms, i.e. of long programs. In the above example, it might be considered somewhat extravagant to declare separate procedures instead of merely substituting the refined program texts for the identifiers. Nevertheless, the gain in clarity of program structure often recommends the use of explicit procedures even in such a simple case. But of course procedures become particularly useful, if the same procedure is to be invoked at several points of the program.

A procedure declaration consists of the symbol PROCEDURE followed by the identifier (together they form the procedure heading), followed by the symbol BEGIN and the statements for which the procedure identifier stands and which are therefore called the *procedure body*. The declaration is terminated by the symbol END and the repetition of the identifier. The latter enables a compiler to detect mismatched endings of statements and declarations. The general syntax of procedure declarations will be given later. A simple example is the following, which computes the sum of a[0] ... a[N-1].

```
PROCEDURE Add;
BEGIN sum := 0.0;
   FOR i := 0 TO N-1 DO sum := a[i] + sum END
END Add
```

The procedure concept becomes even much more useful due to two additional features that are coupled with it, namely the concepts of parameters and of locality of names. Parameters make it possible to invoke the same procedures at different points of the program applying the procedure to different values and variables as determined at the point of the call. Locality of names and objects is a concept considerably enhancing the procedure's role in structuring a program and compartmentalizing its parts. We shall first discuss the concept of locality.

Summarizing, we repeat the following essential points:

1. The procedure aids in exhibiting the inherent structure of a program and in decomposing a programming task.
2. If a procedure is called from two or more points, it reduces the length of the program and therefore the programming task and the potential for programming error. A further economic advantage is the reduction in the size of the compiled code.

11. The Concept of Locality

If we inspect the preceding example of procedure *Add*, we notice that the role of the variable i is strictly confined to the procedure body. This inherent locality should be expressed explicitly, and can be done by declaring i inside the procedure declaration. i thereby becomes a local variable.

```
PROCEDURE Add;
   VAR i: INTEGER;
BEGIN sum := 0.0;
   FOR i := 0 TO N-1 DO sum := a[i] + sum END
END Add
```

In some sense, the procedure declaration assumes the form of a separate program. In fact, any declaration possible in a program, such as constant-, type-, variable-, or procedure declaration, may also occur in a procedure declaration. This implies that procedure declarations may be nested and are defined recursively.

It is good programming practice to declare objects local, i.e. to confine the existence of an object to that procedure in which it has meaning. The procedure, i.e. the section of program text in which a name is declared, is called its *scope*. Since declarations can be nested, scopes are also nestable. The possibility of having objects local to some scope has several consequences. For example, the same name can be used to denote different objects. In fact, this is a most useful consequence, because a programmer is free to choose local identifiers without knowledge of those existing in the surrounding scope, as long as they do not denote objects used in the local scope (in which case he must obviously be aware of them anyhow). This decoupling of knowledge about different program parts is particularly useful and perhaps even vital in the case of large programs.

The rules of scope (validity of identifiers) are as follows:

1. The scope of an identifier is the procedure in which its declaration occurs, and all procedures enclosed by that procedure, subject to rule 2.
2. If an identifier i declared in a procedure P is redeclared in some inner procedure Q enclosed in P, then procedure Q and all procedures enclosed in Q are excluded from the scope of i declared in P.
3. The standard identifiers of Oberon are considered to be declared in an imaginary scope enclosing the program.

These rules may also be remembered by the algorithm in which the declaration of a given identifier i is searched: First, search the declarations of the procedure P in whose body i occurs; if the declaration of i is not among them, continue the search in the procedure surrounding P; then repeat this same rule until the declaration is encountered.

```
VAR a: INTEGER;

PROCEDURE P;
   VAR b: INTEGER;

   PROCEDURE Q;
      VAR b, c: BOOLEAN;
   BEGIN
      (*here a, b (BOOLEAN), c are visible*)
   END Q;

BEGIN
   (*here a, b (INTEGER) are visible*)
END P
```

A consequence of the locality concept and of the rule that a variable does not exist outside its scope is that its value is lost when its declaring procedure is terminated. This implies that, when the same procedure is later called again, this value is unknown. The values of local variables are undefined when the procedure is (re)entered. Hence, if a variable must retain its value between two calls, it must be declared outside the procedure. The "lifetime" of a variable is the time during which its declaring procedure is active.

The use of local declarations has three significant advantages:

1. It makes clear that an object is confined to a procedure, usually a small part of the entire program.

2. It ensures that inadvertent use of a local object by other parts of the program is detected by the compiler.

3. It enables the implementation to minimize storage, because a variable's storage is released when the procedure to which the variable is local is terminated. This storage can then be reused for other variables.

12. Parameters

Procedures may be given parameters. They are the essential feature that make procedures so useful. Consider again the previous example of procedure *Add*. Very likely a program contains several arrays to which the procedure should be applicable. Redefining it for each such array would be cumbersome, inelegant, and can be avoided by introducing its operand as a parameter as follows.

```
PROCEDURE Add(VAR x: Vector);
   VAR i: INTEGER;
BEGIN sum := 0;
   FOR i := 0 TO N-1 DO sum := x[i] + sum END
END Add
```

The parameter x is introduced in the parameter list in the procedure heading. It thereby automatically becomes a local object, in fact is a place-holder for the actual array that is specified in the procedure calls

```
Add(a); ... ; Add(b)
```

The arrays a and b are called *actual parameters* which are substituted for x, which is called *formal parameter*. The formal parameter's specification must contain its type. This enables a compiler to check whether or not an appropriate actual parameter is supplied. We say that a and b, the actual parameters, must be compatible with the formal parameter x. In the example above, its type is Vector, presumably declared in the environment of *Add* as

```
TYPE Vector = ARRAY N OF REAL;
VAR a, b: Vector
```

A better version of Add would include not only the array, but also the result *sum* as a parameter. We shall later return to this example. But first we need explain that there exist two kinds of formal parameters, namely variable and value parameters. The former are characterized by the symbol VAR, the latter by its absence.

We conclude this chapter by giving the syntax of procedure declarations and procedure calls:

```
ProcedureDeclaration = ProcedureHeading ";" block identifier.
ProcedureHeading = "PROCEDURE" identifier [FormalParameters].
block = {declaration} ["BEGIN" StatementSequence] "END".
FormalParameters = "(" [FPSection {";" FPSection}] ")" [":" qualident].
FPSection = ["VAR"] IdentList ":" FormalType.
FormalType = ["ARRAY" "OF"] qualident.

ProcedureCall = designator [ActualParameters].
ActualParameters = "(" [ExpList] ")".
```

Apart from the declaration of modules, we now have also encountered all forms of declarations.

```
declaration = "CONST" {ConstantDeclaration ";"} |
   "TYPE"  {TypeDeclaration ";"} |
   "VAR"  {VariableDeclaration ";"} |
   ProcedureDeclaration ";".
```

12.1 Variable Parameters

As its name indicates, the actual parameter corresponding to a formal variable parameter (specified by the symbol VAR) must be a variable. The formal identifier then stands for that variable.

```
PROCEDURE exchange(VAR x, y: INTEGER);
   VAR z: INTEGER;
```

```
        BEGIN z := x; x := y; y := z
        END exchange
```

The procedure calls

        exchange(a, b); exchange(A[i], A[i+1])

then have the effect of the above three assignments, each with appropriate substitutions made upon the call. The following points should be remembered:

1. Variable parameters may serve to transmit a computed result outside the procedure.
2. The formal parameter acts as a place-holder for the substituted actual parameter.
3. The actual parameter cannot be an expression, and therefore not a constant either, even if no assignment to its formal correspondent is made.
4. If the actual parameter involves indices, these are evaluated when the formal-actual substitution is made.
5. The types of corresponding formal and actual parameters must be the same.

12.2 Value Parameters

Value parameters serve to pass a value from the calling side into the procedure and constitute the predominant case of parameters. The corresponding actual parameter is an expression, of which a variable or a constant are a particular and simple case. The formal value parameter must be considered as a local variable of the indicated type. Upon call, the actual expression is evaluated and the result is assigned to that local variable. As a consequence, the formal parameter may later be assigned new values without affecting any part of the expression. In a sense, actual expression and formal parameter become decoupled as soon as the procedure is entered. As an illustration, we formulate the previously shown program to compute the power $z = x^i$ as a procedure.

```
        PROCEDURE ComputePower(VAR z: REAL; x: REAL; i: INTEGER);
        BEGIN z := 1.0;
           WHILE i > 0 DO
              IF ODD(i) THEN z := z*x END;
              x := x*x; i := i DIV 2
           END
        END ComputePower
```

Possible calls are, for example

        ComputePower(u, 2.5, 3)   for $u := 2.5^3$
        ComputePower(A[i], B[i], 2)   for $A_i := B_i^2$

Concerning value parameters, we must keep in mind that since the formal parameter represents a local variable, storage is needed for that variable. This may be of concern, if its type is an array of many elements. In this case, it is recommended to specify a variable parameter, even if this parameter is used for the import of values only.

Note that in the example above z and x are declared in distinct FPSections, because "VAR z, x: REAL" would classify x as a variable parameter too, making it impossible to place a general expression in its corresponding actual parameter place.

12.3 Open Array Parameters

If a formal parameter type denotes an array structure, its corresponding actual parameter must be an array of the identical type. This implies that it must have elements of identical type and the same bounds of the index range. Often this restriction is rather severe, and more flexibility is highly desirable. It is provided by the facility of the so-called *open array* which requires that the types of the elements of the formal and actual arrays be the same, but leaves the index range of the formal array open. In this case, arrays of any size may be substituted as actual parameters. An open array is specified by the element type preceded by "ARRAY OF". For example, a procedure declared as

        PROCEDURE P(s: ARRAY OF CHAR)

allows calls with character arrays of arbitrary length. The actual length is obtained by calling the standard function LEN(s).

13. Function Procedures

So far we have encountered two possibilities to pass a result from a procedure body to its calling place: the result is either assigned to a non-local variable or to a variable parameter. There exists a third method: the function procedure. It permits the use of the computed result (as an intermediate value) in an expression. The function procedure identifier stands for a computation as well as for the computed result. The procedure declaration is characterized by the indication of the type of the result following the parameter list. As an example, we rephrase the power computation given above.

```
PROCEDURE power(x: REAL; i: INTEGER): REAL;
  VAR z: REAL;
BEGIN z := 1.0;
  WHILE i > 0 DO
    IF ODD(i) THEN z := z*x END;
    x := x*x; i := i DIV 2
  END;
  RETURN z
END power
```

Possible calls are

```
u := power(2.5, 3)
A[i] := power(B[i], 2)
u := x + power(y, i+1) / power(z, i-1)
```

The statement which passes the result consists of the symbol RETURN followed by the expression specifying the result. Return statements may occur at several places in the procedure body and cause termination of its execution. It is good practice, however, that a single return statement is placed immediately preceding the closing symbol END.

Return statements may also be used within normal procedures, in which case no expression follows the RETURN symbol. This facility may serve to signal unusual termination. Such a return statement is implied in the end of every procedure.

Calls inside an expression are called *function designators*. Their syntax is the same as that of procedure calls. However, a parameter list is mandatory, although it may be empty.

```
ReturnStatement = "RETURN" [expression].
```

We now revisit the previous example of adding the elements of an array and formulate it as a function procedure.

```
PROCEDURE sum(VAR a: Vector; n: INTEGER): REAL;
  VAR i: INTEGER; s: REAL;
BEGIN s := 0.0;
  FOR i := 0 TO n-1 DO s := a[i] + s END ;
  RETURN s
END sum
```

This procedure, as previously specified, sums the elements a[0] ... a[n-1], where n is given as a value parameter, and may be different from (but not larger than!) the number of elements N. A more elegant solution specifies *a* as an open array, omitting the explicit indication of the array's size.

```
PROCEDURE sum(VAR x: ARRAY OF REAL): REAL;
  VAR i: INTEGER; s: REAL;
BEGIN s := 0.0;
  FOR i := 0 TO LEN(x)-1 DO s :=  x[i] + s END ;
  RETURN s
END sum
```

Obviously, procedures are capable of generating more than one result by making assignments to several variables. Only one value, however, can be returned as the result of a function. This value, moreover, cannot be of a structured type. Therefore, the other results must be passed to the caller via VAR parameter or assignment to variables that are not local to the function procedure. Consider for example the following procedure which computes a primary result denoted as the function's value and a secondary result used to count the number of times the procedure is called.

```
PROCEDURE square(x: INTEGER): INTEGER;
BEGIN INC(n); RETURN x*x
END square
```

There is nothing remarkable about this example as long as the secondary result is used for its indicated purpose. However, it might be misused as follows:

```
m := square(m) + n
```

Here the secondary result occurs as an argument of the expression containing the function designator itself. The consequence is that e.g. the values

```
square(m) + n   and   n + square(m)
```

differ, seemingly defying the basic law of commutativity of addition.

Assignments of values from within function procedures to non-local variables are called *side-effects*. The programmer should be fully aware of their capability of producing unexpected results when the function is used inappropriately. We summarize:

1. A function procedure specifies a result which is used at its place of call as an argument of an expression.
2. The result of a function procedure cannot be structured.
3. If a function procedure generates secondary results, it is said to have side-effects. These must be used with care. It is advisable to use a regular procedure instead, which passes its results via VAR parameters.
4. We recommend to choose function identifiers which are nouns rather than verbs. The noun then denotes the function's result. Boolean functions are appropriately labelled by an adjective. In contrast, regular procedures should be designated by a verb describing their action.

14. Recursion

Procedures may not only be called, but can call procedures themselves. Since any procedure that is visible can be called, a procedure may call itself. This self-reactivation is called *recursion*. Its use is appropriate when an algorithm is recursively defined and in particular when applied to a recursively defined data structure.

Consider as an example the task of listing all possible permutations of n distinct objects a[0] ... a[n-1]. Calling this operation *Permute(n)*, we can formulate its algorithm as follows:

First, keep a[n-1] in its place and generate all permutations of a[0] ... a[n-2] by calling *Permute(n-1)*, then repeat the same process after having exchanged a[n-1] with a[0]. Continue repeating for all values i = 1 ... n-2. This recipe is formulated as a program as follows, using characters as permuted objects.

```
MODULE Permute;
  IMPORT Texts, Oberon;
  VAR a: ARRAY 20 OF CHAR;
    W: Texts.Writer;

  PROCEDURE permute(k: INTEGER);
    VAR i: INTEGER; t: CHAR;
  BEGIN
    IF k = 0 THEN Texts.WriteString(W, a)
    ELSE permute(k-1);
      FOR i := 0 TO k-1 DO
        t := a[i]; a[i] := a[k]; a[k] := t;
        permute(k-1);
        t := a[i]; a[i] := a[k]; a[k] := t
      END
    END
  END permute;

  PROCEDURE Go*;   (*command*)
    VAR n: INTEGER; S: Texts.Scanner;
  BEGIN Texts.OpenScanner(S, Oberon.Par.text, Oberon.Par.pos); Texts.Scan(S);
```

```
      COPY(S.s, a); n := 0;
      WHILE a[n] > 0X DO INC(n) END ;
      permute(n-1);
      Texts.WriteLn(W); Texts.Append(Oberon.Log, W.buf)
    END Go;

  BEGIN Texts.OpenWriter(W)
  END Permute.
```

The data generated from the command *Permute.Go ABC* are

> ABC BAC CBA BCA ACB CAB

Every chain of recursive calls must terminate at some time, and hence every recursive procedure must place the recursive call within a conditional statement. In the example given above, the recursion terminates when the number of the objects to be permuted is 1. (Concerning input and output conventions, see Ch. 22.3 and 22.4).

The number of possible permutations can easily be derived from the algorithm's recursive definition. We express this number appropriately as a function np(n). Given n elements, there are n choices for the element a[n-1], and with each fixed a[n-1] we obtain np(n-1) permutations. Hence the total number np(n) = n*np(n-1). Evidently np(1) = 1. The computation of np is now expressible as a recursive function procedure.

```
      PROCEDURE np(n: INTEGER): INTEGER;
      BEGIN
        IF n <= 1 THEN RETURN 1
        ELSE RETURN n * np(n-1)
        END
      END np
```

We recognize np as the factorial function, defined as

> $f(n) = 1 \times 2 \times 3 \times \ ... \ \times n$

This formula suggests to program the algorithm using repetition instead of recursion

```
      PROCEDURE np(n: INTEGER): INTEGER;
        VAR p: INTEGER;
      BEGIN p := 1;
        WHILE n > 1 DO p := n*p; DEC(n) END ;
        RETURN p
      END np
```

This formulation will compute the result more efficiently than the recursive version. The reason is that every call requires some "administrative" instructions whose execution costs time. The instructions representing repetition are less time-consuming. Although the difference may not be too relevant, it is recommended to employ a repetitive formulation in place of the recursive one, whenever this is easily possible. It is always possible in principle; however, the repetitive version may complicate and obscure the algorithm to such a degree that the advantages turn into disadvantages. For example, a repetitive form of the procedure permute is much less simple and obvious than the one shown above. To illustrate the usefulness of recursion, two additional examples follow. They typically stem from problems whose solution is naturally found and explained using recursion.

The first example belongs to the class of algorithms which operate on data whose structure is also defined recursively. The specific problem consists of converting simple expressions into their corresponding postfix form, i.e. a form in which the operator follows its operands. An expression shall here be defined using EBNF as follows:

```
      expression = term {("+"|"-") term}.
      term =      factor {("*"|"/") factor}.
      factor =    letter | "(" expression")"|"[" expression "]".
```

Denoting terms as T0, T1, and factors as F0, F1, the rules of conversion are

> T0 + T1   →   T0 T1 +
> T0 - T1   →   T0 T1 -

```
F0 * F1    →   F0 F1 *
F0 / F1    →   F0 F1 /
(E)        →   E
[E]        →   E
```

The following program truthfully mirrors the structure of the syntax of accepted expressions. As the syntax is recursive, so is the program. This close mirroring is the best guarantee for the program's correctness. Note also that similarly repetition in the syntax, expressed by the curly brackets, yields repetition in the program, expressed by while statements. No procedure in this program calls itself directly. Instead, recursion occurs indirectly by a call of *expression* on *term*, which calls *factor*, which calls *expression*. Indirect recursion is obviously much less visible than direct recursion.

This example also illustrates a case of local procedures. Notably, *factor* is declared local to *term*, and *term* local to *expression*, following the rule that objects should preferrably be declared local to the scope in which they are used. This rule may not only be advisable, but even crucial, as demonstrated by the variables *addop* (local to term) and *mulop* (local to factor). If these variables were declared globally, the program would fail. To find the explanation, we must recall the rule that local variables exist (and are given storage) during the time in which their procedure is active. An immediate consequence is that in the case of a recursive call, new incarnations of the local variables are created. Hence, there exist as many as there are levels of recursion. This also implies that a programmer must ensure that the depth of recursion never becomes inordinately large.

```
MODULE Postfix;
  IMPORT Texts, Oberon;
  VAR ch: CHAR;
    W: Texts.Writer; R: Texts.Reader;

PROCEDURE expression;
  VAR addop: CHAR;

  PROCEDURE term;
    VAR mulop: CHAR;

    PROCEDURE factor;
    BEGIN
      IF ch = "(" THEN
          Texts.Read(R, ch); expression;
          WHILE ch # ")" DO Texts.Read(R, ch) END
        ELSIF ch = "[" THEN
          Texts.Read(R, ch); expression;
          WHILE ch # "]" DO Texts.Read(R, ch) END
        ELSE
          WHILE (ch < "a") OR (ch > "z") DO Texts.Read(R, ch) END ;
          Texts.Write(W, ch)
        END ;
        Texts.Read(R, ch)
    END factor;

  BEGIN (*term*) factor;
    WHILE (ch = "*") OR (ch = "/") DO
      mulop := ch; Texts.Read(R, ch); factor; Texts.Write(W, mulop)
    END
  END term;

BEGIN (*expression*) term;
  WHILE (ch = "+") OR (ch = "-") DO
    addop := ch; Texts.Read(R, ch); term; Texts.Write(W, addop)
  END
END expression;

PROCEDURE Parse*;
BEGIN Texts.OpenReader(R, Oberon.Par.text, Oberon.Par.pos); Texts.Read(R, ch);
  WHILE ch # "~" DO
    expression; Texts.WriteLn(W); Texts.Append(Oberon.Log, W.buf);
```

```
      Texts.Read(R, ch)
    END
  END Parse;

  BEGIN Texts.OpenWriter(W);
  END Postfix.
```

A sample of data processed and generated by the command *Postfix.Parse* is shown below.

```
a+b                  ab+
a*b+c                ab*c+
a+b*c                abc*+
a*(b/[c-d])          abcd-/*
```

The next program example demonstrating recursion belongs to the class of problems that search for a solution by trying and testing. A partial "solution" which is once "posted" may, after testing had shown its invalidity, have to be retracted. This kind of approach is therefore also called *backtracking*. Recursion is often very convenient for the formulation of such algorithms.

Our specific example is supposed to find all possible placement of 8 queens on a chess board in such a fashion that none is checking any other piece, i.e. each row, column, and diagonal must contain at most one piece. The approach consists of trying to place a queen in column i, starting with i = 0, and, proceeding to the right, assuming that each column to the left contains a correctly placed queen already. If no place is free in column i, the next column to its left has to be reconsidered. The information necessary to deduce whether or not a given square is still free, is represented by the three global variables called *row, d1, d2* such that

row[j] & d1[i+j] & d2[N-1+i-j] = "the square in column i and row j is free"

Recursion occurs directly in procedure *TryCol*. The auxiliary procedures *PlaceQueen* and *RemoveQueen* could in principle be declared local to *TryCol*. However, there exists a single chess board only (represented by row, d1, d2), and these procedures are appropriately considered as belonging to these global data, and hence not as local to (each incarnation of) *TryCol*.

```
  MODULE Queens;
   IMPORT Texts, Oberon;
   CONST N = 8;  (*no. of rows and columns*)

   VAR x: ARRAY N OF INTEGER;
       (*x[i] = j means: "a queen is on field j in column i"*)
    row: ARRAY N OF BOOLEAN;
       (*row[i] = "no queen on i-th row"*)
    d1:  ARRAY 2*N-1 OF BOOLEAN;
       (*d1[i] = "no queen on i-th upleft to lowright diagonal"*)
    d2:  ARRAY 2*N-1 OF BOOLEAN;
       (*d2[i] = "no queen on i-th lowleft to upright diagonal"*)
    W: Texts.Writer;

   PROCEDURE Clear;
    VAR i: INTEGER;
   BEGIN
    FOR i := 0 TO N-1 DO row[i] := TRUE END ;
    FOR i := 0 TO 2*(N-1) DO d1[i] := TRUE; d2[i] := TRUE END
   END Clear;

   PROCEDURE WriteSolution;
     VAR i: INTEGER;
   BEGIN
    FOR i := 0 TO N-1 DO Texts.WriteInt(W, x[i], 4) END ;
    Texts.WriteLn(W)
   END WriteSolution;

   PROCEDURE PlaceQueen(i, j: INTEGER);
   BEGIN
```

```
    x[i] := j; row[i] := FALSE; d1[i+j] := FALSE; d2[N-1+i-j] := FALSE
  END PlaceQueen;

  PROCEDURE RemoveQueen(i, j: INTEGER);
  BEGIN
    row[i] := TRUE; d1[i+j] := TRUE; d2[N-1+i-j] := TRUE
  END RemoveQueen;

  PROCEDURE TryCol(i: INTEGER);
    VAR j: INTEGER;
  BEGIN
    FOR j := 0 TO N-1 DO
      IF row[j] & d1[i+j] & d2[N-1+i-j] THEN
        PlaceQueen(i, j);
        IF i < N-1 THEN TryCol(i+1) ELSE WriteSolution END ;
        RemoveQueen(i, j)
      END
    END
  END TryCol;

  PROCEDURE Find*;  (*command*)
  BEGIN Clear; TryCol(0); Texts.Append(Oberon.Log, W.buf)
  END Find;

BEGIN Texts.OpenWriter(W)
END Queens.
```

15. Type Declarations

Every variable declaration specifies the variable's type as its constant property. The type can be one of the standard, primitive types, or it may be of a type declared in the program itself. Type declarations have the form

TypeDeclaration = identifier "=" type.

They are preceded by the symbol TYPE. Types are classified into unstructured and structured types. Each type essentially defines the set of values which a variable of this type may assume. A value of an unstructured type is an atomic unit, whereas a value of structured type has components (elements). For example, the type INTEGER is unstructured; its elements are atomic. It does not make sense, e.g. to refer to the third bit of the value 13; the circumstance that a number may "have a third bit", or a second digit, is a characteristic of its (internal) representation, which intentionally is to remain unknown.

In the following sections we shall show how to declare structured types. We distinguish between various structuring methods of which we have so far encountered the array only. In addition, there exists the record type. A facility to introduce structures that vary dynamically during program execution is based on the concept of pointers and will be discussed in a separate chapter.

type = qualident | ArrayType | RecordType | PointerType | ProcedureType.

Before proceeding to the various kinds of types, we note that in general, if a type T is declared by the declaration

TYPE T = someType

and a variable t is declared as

VAR t: T

then these two declarations can always be merged into the single declaration

VAR t: someType

However, in this case t's type has no explicit name and therefore remains anonymous. Typically, record types are given explicit names.

The concept of type is important, because it divides a program's set of variables into disjoint classes. Inadvertent assignments among members of different classes can therefore be detected by a mere inspection of the program text without executing the program. Given, for example, the declarations

    VAR b: BOOLEAN; i: INTEGER; x: REAL

the assignment b := i is impossible, because the types of b and i are incompatible. Two types are said to be compatible, if they are declared as equal or satisfy certain compatibility rules: If one type T0 is included in another type T1, then the assignment t1 := t0 is admissible, but not t0 := t1. Type inclusion is restricted to the standard numerical types. We note that x := i is admisible, whereas i := x is not.

    SHORTINT $\subseteq$ INTEGER $\subseteq$ LONGINT $\subseteq$ REAL $\subseteq$ LONGREAL

16. Record Types

In an array all elements are of the same type. In contrast to the array, the record structure offers the possibility to declare a collection of elements as a unit even if the elements are of different types. The following examples are typical cases where the record is the appropriate choice of structuring method. A date consists of three elements, namely day, month, and year. A description of a person may consist of the person's names, sex, identification number, and birthdate. This is expressed by the following type declarations

```
Date  =    RECORD day, month, year: INTEGER END

Person  =  RECORD
                firstName, lastName: ARRAY 32 OF CHAR;
                male: BOOLEAN;
                idno: INTEGER;
                birth: Date
           END
```

The record structure makes it possible to refer either to the entire collection of data or to individual elements. Elements of a record are also called *record fields*, and their names are called *field identifiers*. This stems from the habit of looking at such data as forms or tables drawn on paper with individual fields delineated as rectangles and labelled with field names. Similar to arrays, where we denote the i-th element of an array *a* by *a[i]*, i.e. by the array identifier followed by an index, we denote the field f of a record *r* by *r.f* , i.e. by the record identifier followed by the field's name. For example, given the variables

```
d1, d2 :  Date;
p1, p2 :  Person;
student:  ARRAY 100 OF Person
```

we can construct, for example, the following variable designators:

```
d1.day
d2.month
p1.firstName
p1.birth
```

These examples show that fields may themselves be structured. Similarly, records may be elements of array or record structures, i.e. there exists the possibility to construct hierarchies of structures. As a consequence, selectors of elements can be sequenced, as shown by the following examples of designators. The multi-dimensional array discussed in the chapter on arrays now appears as a particular case of these structuring hierarchies.

```
p1.lastName[0]
p2.birth.year
student[23].idno
student[k].firstName[0]
```

At first sight the record may appear as a generalized array, because it relaxes the restriction that all elements be of the same type. However, in another aspect it is more restrictive than the array: the selector of the element must be a fixed field identifier, whereas the index selecting an array element may be an expression, i.e. a result of previous computations.

It is important to note that a record may assume arbitrary combinations of its field's values. Hence, in the example of the type Date, a value *day* = 31 may coexist with *month* = 2, although this is not an actual date.

The syntax of a record declaration is defined as follows.

        RecordType = "RECORD" FieldListSequence "END".
        FieldListSequence = FieldList {";" FieldList}.
        FieldList = [IdentList ":" type].

and that of a designator is

        designator = qualident {selector} .
        selector = "." identifier | "[" ExpList "]" | "^".

Note: Oberon's concept of record *type extension* is not discussed in this book.

17. Dynamic Data Structures and Pointers

Array and record structures share the common property that they are static. This implies that variables of such a structure maintain the same structure during the whole time of their existence. In many applications, this is an intolerable restriction; they require data which do not only change their value, but also their composition, size, and structure. Typical examples are lists and trees that grow and shrink dynamically. Instead of providing list and tree structures, a collection that for some applications would again not suffice, Oberon offers a basic tool to construct arbitrary structures. This is the *pointer type*.

Every complex data structure ultimately consists of elements whose structure is static. Pointers, i.e. values of pointer types, are themselves not structured, but rather are used to establish relationships among those static elements, usually called *nodes*. We also say that pointers *link* elements or *point to* elements. Evidently, different pointer variables may point to the same element, hence providing the possibility to compose arbitrarily complex structures, and at the same time opening many pitfalls for programming mistakes that are difficult to pinpoint. Operating with pointers indeed requires utmost care.

Pointers in Oberon cannot point to arbitrary variables. The type of variable to which they point must be specified in the pointer type's declaration, and the pointer type is said to be *bound* to the referenced object's type. Example:

        TYPE Node = POINTER TO NodeDescriptor
        VAR p0, p1: Node

Here *Node* (and thereby also variables p0 and p1) are bound to the type *NodeDescriptor*, i.e. they can point to variables of type NodeDescriptor only. These variables, typically of record type, are *not* created by the declaration of p0 and p1. Instead, they are created by a call an *allocation procedure*. In Oberon, it is represented by a predefined operator called NEW. The statement NEW(p0) then creates a (record) variable of type NodeDescriptor and assigns a pointer referring to that variable (i.e. a value of type Node) to p0. The created variable is said to be dynamically created (allocated); it has no name, is anonymous, and can be accessed only via a pointer using the dereferencing operator ^. The said variable is denoted by the designator p0^. If the referenced variable is of record type with, say, fields x and y, then the fileds are denoted, for example, by p0^.x, p1^.y. Since in this case it is clear that field selection applies to the referenced record rather than to the pointer, the designators may be abbreviated into p0.x and p1.y.

        PointerType = "POINTER" "TO" type.

What really makes pointers such a powerful tool is the circumstance that they may point to variables which themselves contain pointers. This reminds us of procedures that call procedures and thereby introduce recursion. In fact, pointers are the tool to implement recursively defined data structures (such as lists and trees). The nature of the recursive data structure is evident from the declaration of the type of its elements.

Just as every recursion of procedure activation must terminate at some time, so must every recursion in referencing terminate at some point. The role of the if statement to terminate procedural recursion is here taken by the special pointer value NIL terminating referencing recursion. NIL points to no object. It is an obvious consequence that a designator of the form p^ (p^.x, p.x) must never be evaluated, if p = NIL. We summarize the following essential points.

1. Every pointer type is bound to a type; its values are pointers which point to variables of that type.
2. The referenced variables are anonymous and can be accessed via pointers only.
3. The referenced variables are dynamically created by an allocation procedure which assigns the variable's pointer to p.

4. The pointer constant NIL belongs to every pointer type and points to no object.

5. The variable referenced by a pointer p is denoted by the designator p^. In order for p^ to be meaningful, p must not have the value NIL.

6. Field designators of the form p^.x may be abbreviated to p.x.

Lists, also called linear lists or chains, are characterized by consisting of record typed nodes that each have exactly one element being a pointer to a record of the same type as itself. This implies recursion. A list pointer declaration then assumes the characteristic form

```
List  =  POINTER TO ListDesc

ListDesc  =
 RECORD
    key:  INTEGER;
    Data: ...
    next: List
 END
```

"Data" actually stands for any number of fields representing data pertaining to the listed node. *Key* is part of these data; it is mentioned separately here because it is quite common to associate with each element a unique identifying key, and also because it will be used in subsequent examples of operations on lists. The essential ingredient here, however, is the field *next*, so labelled because it evidently is the pointer to the next element in the list. Direct recursion in data type declarations is not permitted for the obvious reason that there would be no evident termination. The declaration given above cannot be abbreviated into

```
List  =
 RECORD
    key: INTEGER; next: List
 END
```

Assume now that a list is accessible in a program via its first element, denoted by the pointer variable

```
first:  List
```

The empty list is represented by first = NIL. A non-empty list is most conveniently constructed by inserting new elements at its front. The following assignments are needed to insert one element (let p be an auxiliary variable of type *List*)

```
NEW(p); (* assign values to p.key and p.data *)  p.next := first; first := p
```

Having constructed a list by repeated insertion of nodes, we may wish to search the list for a node with key value equal to a given x. We evidently use a repetition; the while statement is appropriate, because we do not know the number of nodes (and hence repetitions) beforehand. It is wise to include the case of the empty list!

```
p := first;
WHILE (p # NIL) & (p.key # x) DO p := p.next ND ;
IF p # NIL THEN found END
```

We draw attention to the fact that here we make use of the rule that the term b is not evaluated, if in the expression a & b the factor a is found to be FALSE. If this rule would not hold, the factor p.key # x might be evaluated with p = NIL, which is illegal.

The second frequently encountered dynamic data structure is the *tree*. It is characterized by its nodes having *n* pointer fields each, where n is the *degree* of the tree. The common and in some sense optimal case is the binary tree with n = 2. Lists now appear as degenerate trees of degree 1. The respective declarations are

```
Tree  = POINTER TO TreeNode;
TreeNode =
 RECORD
    key: INTEGER;
    data: ...
    left, right: Tree
 END
```

The place of the variable *first* in the case of lists is taken by a variable to be

    root: Tree

with root = NIL denoting the empty tree. Trees are commonly used to represent collections of data in order of ascending key values, making retrieval very efficient. The following statements represent a search in an ordered binary tree, whose similarity to the binary search in an ordered array is remarkable. Again, p is an auxiliary variable (of type Tree).

```
p := root;
WHILE (p # NIL) & (p.key # x) DO
  IF p.key < x  THEN  p := p.right
  ELSE p := p.left
  END
END ;
IF p # NIL THEN found END
```

This example is a repetitive version of the tree search. Next we show the recursive version. It is, in addition, extended such that a new node is created and inserted at the appropriate place, whenever no node with key value x exists.

```
PROCEDURE search(VAR p: Tree; x: INTEGER): Tree;
  VAR q: Tree;
BEGIN
  IF p # NIL THEN
    IF p.key < x THEN
      q := search (p.right, x)
    ELSIF p.key > x THEN
      q := search (p^.left, x)
    ELSE
      q := p
    END
  ELSE (*not found, hence insert*)
    NEW(q); q.key := x;  q.left := NIL;  q.right := NIL
  END ;
  RETURN q
END search
```

The call *search(root, x)* now stands for a search of x in the tree represented by the variable *root*.

And this concludes our examples of operations on lists and trees to illustrate pointer handling. Lists and trees have nodes which are all of the same type. We draw attention to the fact that the pointer facility admits the construction of even more general data structures consisting of nodes of various types. Typical for all these structures is that all nodes are declared as record types. Hence, the record emerges as a particularly useful data structure in conjunction with pointers.

Creation of nodes is expressed by the standard procedure NEW, which is part of a system's storage management. We assume that retrieval of storage is performed automatically by a so-called storage reclamation mechanism, also called *garbage collector*. It relies on the fact that records that are no longer referenced by any pointer may be recollected, i.e. their storage can be recycled.

18. Procedure Types

So far, we have regarded procedures exclusively as program parts, i.e. as texts that specify actions to be performed on variables whose values are numbers, logical values, characters, etc. However, we may take the view that procedures themselves are objects that can be assigned to variables. In this light, a procedure declaration appears as a special kind of constant declaration, the value of this constant being a procedure. If we allow variables in addition to constants, it must be possible to declare types whose values are procedures. These are called *procedure types*.

A procedure type declaration specifies the number and the types of parameters and, if it is to be a function procedure, the type of the result. For example, a procedure type with one REAL argument and a result of the same type is declared by

    Function = PROCEDURE (x: REAL): REAL

The mathematical functions of sine, cosine, square root, exponential and logarithm are all of this type. The general syntax is

    ProcedureType = "PROCEDURE" [FormalTypeList].
    FormalTypeList =
      "(" [["VAR"] FormalType {"," ["VAR"] FormalType}] ")" [":" identifier].

If we now declare a variable *f: Function*, the assignment *f := Math.sin* is possible. Subsequently the call *f(x)* is equivalent to *Math.sin(x)*.

It is now also possible to declare functions and procedures which themselves have functions and procedures as parameters. For example, a function procedure to integrate any function over a certain interval can be expressed in the following way. It adds the values of the given parametric function ofver the given interval a to b in n steps.

    PROCEDURE integral (f: Function; a, b: REAL; n: INTEGER): REAL;
      VAR i: INTEGER; sum, dx: REAL;
    BEGIN dx := (b-a)/n; sum := 0;
      FOR i := 0 TO n-1 DO sum := f(a + i*dx + 0.5*dx) + sum END ;
      RETURN sum * dx
    END integral

The integrations of, for example, the sine function over the interval 0 to π, and that of the exponential function from 0 to 1 are expressed simply as

    integral (Math.sin, 0, 3.14159, 20)
    integral (Math.exp, 0, 1, 20)

and they yield 2.0020 and 1.7181 respectively, which approximate the correct values 2*cos(0) and e − 1).

To conclude, we emphasize the restriction of Oberon that procedures that are assigned to variables or are passed as parameters, must not be declared local to any other procedure. Neither can they be standard procedures (such as ODD, INCL, etc.).

Part 3

19. Modules

Modules are the most important feature distinguishing Oberon from its ancestor Pascal. We have already encountered modules, simply because every program is a module. However, most systems are not only a single module, but rather a set of several modules, each module containing declared objects such as constants, variables, procedures, and types. Objects declared in a modules M0 can be referenced in another module M1, if they are explicitly made to be known in M1, i.e. if they are *imported* by M1. In the examples of the preceding chapters, we have typically imported procedures for input and output from modules containing collections of frequently used procedures. These procedures are actually part of our program, even if we have not programmed them and they are textually disjoint.

The key point is that modules can be kept in a program "library" and are automatically referenced when a programmer's program is loaded and executed. Hence it is possible to prepare collections of frequently used operations (such as for input and output), and to avoid reprogramming them each time a program needs such operations. Modern implementations go even one step further and offer what is called *separate compilation*. This signifies that such modules are not stored in the program library in source form as Oberon texts, but rather in compiled form. Upon program loading, the compiled (main) module is joined (linked) with the precompiled modules from which it imports objects. In this case, the compiler must also have access to descriptions of the objects of the previously compiled, imported modules, when the importing program is compiled. This facility distinguishes separate compilation from *independent compilation* as it exists in typical implementations of Fortran, Pascal, and assemblers.

Every subsidiary module may again import objects from other modules. A program therefore constitutes an entire hierarchy of modules. The main program is said to be at the highest level, those modules which do not import objects at all being at the lowest level. Usually, a programmer is not even aware of this hierarchy, because his modules import objects from modules that he has not programmed himself; therefore he is unaware of their imports and of the module hierarchy below them. In principle, however, his program is the text written by himself, extended with the texts of the imported modules.

These extensions are usually quite large (even if the direct imports consist of a few output procedures only). In principle, the indirect imports constitute the entire *environment* or operating system. In a single-user computer system, there is virtually no need for any parts that are neither directly nor indirectly imported by the main program. However, some modules, such as the basic input/output and file system, may be required by all programs and therefore become de facto resident and may therefore be regarded as the operating system.

The principal motivation behind the partitioning of a program into modules is - beside the use of modules provided by other programmers - the establishment of a hierarchy of abstractions. For example, in the previously encountered cases of imported input/output procedures, we merely wish to have them available, but do not need to know - or rather do not wish to bother to learn - how these procedures function in detail. To abstract means to "take away" from the essentials and thereby to ignore certain details. Each module constitutes an abstraction, if we regard it from "the outside". We even wish to go one step further: we wish not only to ignore the details of its innards, but to hide them. The primary reason for this wish is that if the innards are protected from outside access, we can guarantee their correct functioning, thereby being able to limit the area of error search in the case of a malfunctioning program.

The second, but not less important reason is to make it possible to change (improve) the innards of imported modules without having to change (and/or recompile) the importing modules. This effective decoupling of modules is indispensible for the development of large programs, in particular, if modules are developed by different people, and if we regard the operating system as the low section of a program's module hierarchy. Without decoupling, any change or correction in an operating system or in library modules would become virtually impossible.

A direct consequence of this need for decoupling is the necessity for a textual separation of the essentials from the details. The essentials of a module are the properties of objects that are importable into other modules; the details are those parts that are to be hidden and protected. In Oberon, objects

to be visible in other modules are specially marked. The mark is an asterisk following the identifier in its declaration.

An importer of a module needs to know the properties of those objects only which are exported by the imported module. It is customary and convenient to provide an extract of the module containing those relevant declarations only. We call it the *definition* of a module. The definition forms a contract between the client (importer) and the designer of a module. It is therefore also called the *interface* of the module. The implementation part remains the property of the module's designer. As long as the designer alters details without affecting the definition, he need not report his activity to the clients of his module. Modules are compiled separately, and are therefore called *compilation units*.

Concluding this introduction to the module concept, we postulate that adequate implementations provide full type compatibility checking between objects, independent of whether these objects are declared in the same or in different modules, i.e. the checking mechanism of the compiler functions also across module boundaries. However, the programmer must realize that this checking provides no absolute safeguard against mistakes. After all, it concerns formal, syntactic aspects only; it does not cover the semantics. It would not detect, for instance, the replacement of the algorithm for the sine function by that of the cosine function. However, it must not be regarded as the duty of a compiler to protect programmers against malicious "colleagues".

20. Definitions and Implementations

A definition or interface specifies those properties of a module which are relevant to its clients (importers, users). It consists of the declarations of the exported identifiers. We will here consider definitions and implementations as disjoint texts. However, in Oberon there is only one text, namely the module, in which the exported objects are explicitly marked (by an asterisk). The definition may here be considered as an extract of the module text.

Variables declared in a definition are global in the sense that they exist during the entire lifetime of the program, although they are visible and accessible only in those modules (clients) which import the module of their declaration. In other modules they remain invisible. Such variables are typically state variables. But in general, export of variables should be used rarely, and if so, preferrably imported variables should be read only.

In definitions, procedures are specified by their headings only. A heading specifies the procedure's parameters, i.e. their types, and in the case of function procedures the result type. This information constitutes the procedure's *signature*.

If a type is declared in a definition module, the full details of its declaration are visible in importing modules. In the case of a record type, the definition specifies only those fields, which are visisble by clients.

The following simple example exhibits the essential characteristics of modules, although typically modules are considerably larger pieces of program and contain a longer list of declarations. This example exports two procedures, *put* and *get*, adding data to and fetching data from a buffer which is hidden. Effectively, the buffer can only be accessed through these two procedures; it is therefore possible to guarantee the buffer's proper functioning, i.e. sequential access without loss of elements.

```
DEFINITION Buffer;
  VAR nonempty, nonfull: BOOLEAN;
  PROCEDURE put(x: INTEGER);
  PROCEDURE get(VAR x: INTEGER);
END Buffer.
```

This definition part contains all the information about the buffer that a client is supposed to know. The details of its operation, its realization, are contained in the corresponding implementation.

```
MODULE Buffer;  (*cyclic buffer of integers*)
  CONST  N = 100;
  VAR nonempty*, nonfull*: BOOLEAN;
    in, out, n: INTEGER;
    buf: ARRAY N OF INTEGER;

  PROCEDURE put*(x: INTEGER);
  BEGIN
```

```
        IF n < N THEN
           buf[in] := x; in := (in+1) MOD N;
           INC(n); nonfull := n < N; nonempty := TRUE
        END
     END put;

     PROCEDURE get*(VAR x: INTEGER);
     BEGIN
       IF n > 0 THEN
          x := buf[out]; out := (out+1) MOD N;
          DEC(n); nonempty := n > 0; nonfull := TRUE
       END
     END get;

   BEGIN (*initialize*)  n := 0; in := 0; out := 0;
     nonempty := FALSE; nonfull := TRUE
   END Buffer.
```

The definition is the relevant extract from the module losted here, implementing a fifo (first in first out) queue. This fact, however, is not evident from the definition alone; normally the semantics are mentioned in the form of a comment or other documentation. Such comments will usually explain what the module performs, but not how this is achieved. Therefore, different implementations may be provided for the same definition. The differences may lie in the details of the mechanism; for example, the buffer might be represented as a linked list instead of an array (allocating buffer portions as needed, hence not limiting the buffer's size). Or, the differences may even lie in the semantics. The following program implements a stack (i.e. a lifo queue) instead of a fifo queue, nevertheless fitting the same definition part. Any change in a module's semantics necessitates corresponding adjustments in the module's clients, and must therefore be made with utmost care.

```
   MODULE Buffer;  (*stack of integers*)
     CONST  N = 100;
     VAR n: INTEGER;
        nonempty*, nonfull*: BOOLEAN;
        buf: ARRAY N OF INTEGER;

     PROCEDURE put*(x: INTEGER);
     BEGIN
       IF n < N THEN
          buf[n] := x; INC(n); nonfull := n < N; nonempty := TRUE
       END
     END put;

     PROCEDURE get*(VAR x: INTEGER);
     BEGIN
       IF n > 0 THEN
          DEC(n); x := buf[n]; nonempty := n > 0; nonfull := TRUE
       END
     END get;

   BEGIN n := 0; nonempty := FALSE; nonfull := TRUE
   END Buffer.
```

Evidently, *nonempty* is the precondition for *get*, and *nonfull* is the precondition of *put*. This concludes this introductory example.

The syntax of modules and definitions is

```
   Definition =
     "DEFINITION"  identifier ";" {import} {declaration}
     "END" identifier "." .

   declaration = "CONST" {ConstantDeclaration ";"}  |
     "TYPE"  {identifier ["=" type]  ";"}    |
     "VAR"  {VariableDeclaration ";" }     |
     ProcedureHeading ";".
```

Module = "MODULE" identifier ";" ["IMPORT" IdList} block identifier.

The import list names the modules which are imported. If a client module M1 is a client of M0, i.e. imports M0, then the objects exported from M0, say *a, b,* are referenced by qualified identifiers of the form M.x, for example M0.a, M0.b. This facility permits to import different objects with the same name from different modules and to avoid conflicts of names. Standard identifiers are automatically imported into all modules.

The possibility to publicize a module in the form of its definition and at the same time to retain its operational details hidden in its implementation specification, is particularly convenient for the establishment of program libraries. Such collections of standard routines belong to every programming environment. They typically include routines for input and output operations, for file handling, and for the computation of mathematical functions. Although there exists no rigid standard for Oberon, the modules *Files* and *Texts* can be considered as standard modules available in all implementations of Oberon (see Ch. 22). Here we present the definition of module *Math* as a first example, featuring standard mathematical functions:.

```
DEFINITION Math;
  PROCEDURE sqrt(x: REAL): REAL;
  PROCEDURE exp(x: REAL): REAL;
  PROCEDURE ln(x: REAL): REAL;
  PROCEDURE sin(x: REAL): REAL;
  PROCEDURE cos(x: REAL): REAL;
  PROCEDURE arctan(x: REAL): REAL;
  PROCEDURE real(x: INTEGER): REAL;
  PROCEDURE entier(x: REAL): INTEGER;
END Math.
```

21. Program Decomposition into Modules

The quality of a program has many aspects and is an elusive property. A user of a program may judge it according to its efficiency, reliability, or convenience of user dialog. Whereas efficiency can be expressed in terms of numbers, convenience of usage is rather a matter of personal judgement, and all too often a program's usage is called convenient as long as it is conventional. An engineer of a program may judge its quality according to its clarity and perspicuity, again rather elusive and subjective properties. However, if a property cannot be expressed in terms of precise numbers, this is no reason for classifying it as irrelevant. In fact, program clarity is enormously important, and to demonstrate the correctness of a program is ultimately a matter of convincing a person that the program is trustworthy. How can we approach this goal? After all, complicated tasks usually do inherently require complex algorithms, and this implies a myriad of details. And the details are the jungle in which the devil hides.

The only salvation lies in structure. A program must be decomposed into partitions which can be considered one at a time without too much regard for the remaining parts. At the lowest level the elements of the structure are statements, at the next level procedures, and at the highest level modules. In parallel with program structuring proceeds the structuring of data into arrays, records, etc. at the lower levels, and through association of variables with procedures and modules at the higher levels. The essence of programming is finding the right (or at least an appropriate) structure, and the experienced programmer is the person who has the intuition to find it at the stage of initial conception instead of during a gradual process of improvements and modifications. Yet, the programmer who has the courage to restructure when a better solution emerges is still much better off than the one who resigns and elaborates a program on the basis of a clearly inadequate structure, for this leads to those products that no one else, and ultimately not even the originator himself can "understand".

Even if there does not exist a recipe to determine the optimal structure of a program, there have emerged some criteria for guidance in the process of finding good and avoiding bad structure. A basic rule is that decomposition should be such that the connections between partitions are simple or "thin". A perhaps oversimplifying criterion for the thickness of a connection - also called *interface* - between two parts is the number of items that take part in it. Specifically, the interface of two modules is sketched in terms of the module's import lists, and a measure for the interface's thickness is the number of imported items. Hence, we must find a modularization which makes the import lists short.

Naturally, it is difficult to find an optimum, for, the import lists would be shortest, i.e. they would disappear if the entire program were collapsed into a single module: a clearly undesirable solution.

The distinctive property of the module as the largest structuring unit is its capability to hide details and thereby to establish a new level of abstraction. This property is used in several forms; we can distinguish between the following cases:

1. The module separates two kinds of data representation and contains the collection of procedures that perform the data conversion between the two levels. The typical example is a module for conversion of numbers from their abstract, atomic representation into sequences of decimal digits and vice-versa. Such modules contain no data of their own, they are typically packages of procedures.

2. The essence of a module is a set of data. It hides the details of the data representation by granting access to these data through calls of its exported procedures only. An example of this case is a module which contains a data set storing individual items organized in a way that items are found quickly. Another is a module whose hidden data set is a disk store; it hides the peculiar details necessary to operate the disk drive.

3. The module exports a data type and exports its associated operations. Typically, in Oberon such a module exports one or several types in opaque mode (sometimes these are also called private types). It thereby hides the details of the data type's structure and also the details of the operations. By hiding them, it is possible to guarantee the validity of postulated invariant properties of each variable of such a private type. The difference to modules of class 2 is that here variables of the private types are declared in the client modules, whereas in class 2 modules the variable is itself hidden. Typical examples are the queue and stack types, and perhaps the most successful such data abstraction is the sequential file, also known as a *stream*.

This classification is not absolute. It cannot be, because all kinds share the common goal of hiding details. Nevertheless, we can formulate a few rules that serve as guidelines in the design of modules:

1. Keep the number of imported modules small.
2. Keep the number of imported modules in definitions even smaller.
3. Export of variables should be considered as the exception, and imported variables should be treated as "read-only" objects.

We conclude this chapter with an example that essentially falls into category 3. Let our stated goal be the design of a cross reference generator for Oberon programs. More precisely, the program's purpose is to read a text and to generate (1) a listing of the text with added line numbers and (2) a table of all encountered words (identifiers) in alphabetical order, each with a list of the numbers of the lines in which the word occurs. Moreover, comments and strings are to be skipped (i.e. their words are not to be listed), and Oberon key symbols are not to be listed either.

We quickly recognize the task as being divisible into the scanning of the source text (eliminating the parts that are to be skipped and ignored), and the recording and subsequent tabulating of the words. The first part is conveniently performed by the main program module, the latter by a subsidiary module which hides the data set and makes it accessible through two procedures: *Insert* (i.e. include a word) and *List* (i.e. generate the requested table). A third module is used to generate the representation of numbers as sequences of decimal digits. The three principal modules involved are called XREF, TableHandler, and Texts.

We begin by presenting the main program XREF that scans the source text. A binary search is used to recognize key words. The data set is of the type *Table*, imported from the *TableHandler*.

```
DEFINITION TableHandler;
 CONST WordLength;
 TYPE Word;
 PROCEDURE Init (VAR w: Word);

 PROCEDURE Insert(VAR s: ARRAY OF CHAR; ln: INTEGER; VAR w: Word);
   (*enter pair s, ln in structuire w *)

 PROCEDURE List(w: Word)
END TableHandler.
```

```
    MODULE XREF;
      IMPORT Texts, Oberon, TableHandler;
      CONST N = 32;   (* No. of keywords *)

      VAR lno: INTEGER;  (*current line number*)
        Tab: TableHandler.Word;
        W: Texts.Writer;
        key: ARRAY N, 10 OF CHAR;  (*table of key words*)

      PROCEDURE heading;
      BEGIN INC(lno); Texts.WriteInt(W, lno, 5); Texts.Write(W, " ")
      END heading;

    PROCEDURE Scan*;   (*command*)
      VAR beg, end, time: LONGINT;
        k, m, l, r: INTEGER;
        ch: CHAR;
        id: ARRAY TableHandler.WordLength OF CHAR;
        T: Texts.Text; R: Texts.Reader;

      PROCEDURE copy;
      BEGIN Texts.Write(W, ch); Texts.Read(R, ch);
      END copy;

    BEGIN TableHandler.Init(Tab);
      Oberon.GetSelection(T, beg, end, time);
      IF time >= 0 THEN
        lno := 0; heading;
        Texts.OpenReader(R, T, beg); Texts.Read(R, ch);
        REPEAT
          IF (CAP(ch) >= "A") & (CAP(ch) <= "Z") THEN (*word*)
            k := 0;
            REPEAT id[k] := ch; INC(k); copy
            UNTIL ~(("A" <= CAP(ch)) & (CAP(ch) <= "Z") OR ("a" <= ch) & (ch <= "z"));
            id[k] := 0X;
            l := 0; r := N;  (*binary search for key word*)
            REPEAT m := (l+r) DIV 2;
              IF key[m] < id THEN l := m+1 ELSE r := m END ;
            UNTIL l >= r;
            IF (r = N) OR (id # key[r]) THEN TableHandler.Insert(id, lno, Tab) END
          ELSIF (ch >= "0") & (ch <= "9") THEN (*number*)
            REPEAT copy UNTIL ~("0" <= ch) & (ch <= "9")
          ELSIF ch = "(" THEN copy;
            IF ch = "*" THEN  (*skip comment*)
              REPEAT
                REPEAT
                  IF ch = 0DX THEN copy; heading ELSE copy END
                UNTIL ch = "*";
                copy
              UNTIL ch = ")";
              copy
            END
          ELSIF ch = 22X THEN (*string*)
            REPEAT copy UNTIL ch = "”";
            copy
          ELSIF ch = 0DX THEN (*end of line*)
            copy; heading
          ELSE copy
          END
        UNTIL R.eot;
        Texts.WriteLn(W); Texts.Append(Oberon.Log, W.buf);
        TableHandler.List(Tab)
```

```
    END
END Scan;

BEGIN Texts.OpenWriter(W);
  key[ 0] := "AND";        key[ 1] := "ARRAY";        key[ 2] := "BEGIN";
  key[ 3] := "BOOLEAN"; key[ 4] := "BY";        key[ 5] := "CASE";
  key[ 7] := "CONST";      key[ 7] := "DIV";        key[ 8] := "DO";
  key[ 9] := "ELSE";        key[10] := "ELSIF";        key[11] := "END";
  key[12] := "FOR";        key[13] := "IF";        key[14] := "IMPORT";
  key[15] := "IN";        key[16] := "MOD";        key[17] := "MODULE";
  key[18] := "NOT";        key[19] := "OF";        key[20] := "OR";
  key[21] := "POINTER";  key[22] := "PROCEDURE"; key[23] := "RECORD";
  key[24] := "REPEAT";    key[25] := "RETURN";    key[26] := "THEN";
  key[27] := "TO";        key[28] := "TYPE";        key[29] := "UNTIL";
  key[30] := "VAR";        key[31] := "WHILE"
END XREF.
```

Next we present the table handler. As seen from its definition part, it exports the private type *Table* and its operations *Insert* and *List*. Notably the structure of the tables, and thereby also the access and search algorithms remain hidden. The two most likely choices are the organizations of binary trees and of a hash table. Here we opt for the former. The example is therefore a further illustration of the use of pointers and dynamic data structures. The module contains a procedure to search and insert a tree element, and a procedure that traverses the tree for the required tabulation (see also the Chapter on dynamic data structures). Each tree node is a record with fields for the key, the left and right descendants, and (the head of) a list of the line numbers. Another form of representation might be chosen, for example a balanced tree, and such a new implementaion might be provided without clients being aware of the change, because modules can be compiled separately. This is the key of modularization and constructing large systems. It requires, however, that the modules' interfaces are wisely chosen.

```
    MODULE TableHandler;
     IMPORT Texts, Oberon;

     CONST WordLength* = 32;

     TYPE Item =  POINTER TO RECORD
        num: INTEGER; next: Item
      END ;

      Word* =  POINTER TO RECORD
        key: ARRAY WordLength OF CHAR;
        first: Item;      (*list head*)
        left, right: Word
      END ;

     VAR W: Texts.Writer;

     PROCEDURE Init*(VAR w: Word);
     BEGIN w := NIL
     END Init;

     PROCEDURE Insert*(VAR s: ARRAY OF CHAR; lno: INTEGER; VAR w: Word);
      (*search node with name s*)
      VAR w0: Word; t: Item;
     BEGIN
      IF w = NIL THEN (*insert new word*)
        NEW(w0); COPY(s, w0.key); w0.left := NIL; w0.right := NIL;
        NEW(t); t.num := lno; t.next := NIL; w0.first := t; w := w0
      ELSIF s < w.key THEN Insert(s, lno, w.left)
      ELSIF s > w.key THEN Insert(s, lno, w.right)
      ELSE NEW(t); t.num := lno; t.next := w.first; w.first := t
      END
     END Insert;
```

```
    PROCEDURE write(w: Word);
      VAR t: Item;
    BEGIN Texts.WriteString(W, w.key); t := w.first;
      REPEAT Texts.WriteInt(W, t.num, 5); t := t.next UNTIL t = NIL;
      Texts.WriteLn(W)
    END write;

    PROCEDURE List*(w: Word);
    BEGIN
      IF w # NIL THEN List(w.left); write(w); List(w.right) END ;
      Texts.Append(Oberon.Log, W.buf)
    END List;

  BEGIN Texts.OpenWriter(W)
  END TableHandler.
```

22. The Concept of a Sequence

22.1. About input and output

The usefulness and the success of high-level programming languages rests on the principle of abstraction, the hiding of details that pertain to the computer which is used to execute the program rather than to the algorithm expressed by the program. The domain that has most persistently relied on abstraction is that of input and output operations. This is not surprising, because input and output inherently involve the activation of devices that are peripheral to the computer, and whose structure, function, and operation differ strongly among various kinds and brands. Many programming languages have typically incorporated statements for reading and writing data in sequential form without reference to specific devices or storage media. Such an abstraction has many advantages, but there exist always applications where some property of some device is to be utilized that is poorly, if at all, available through the standard statements of the language. Also, generality is usually costly, and consequently operations that are conveniently implemented for some devices may be inefficient if applied to other devices. Hence, there also exists a genuine desire to make transparent the properties of some devices for applications that require their efficient utilization. Simplification and generalization by suppression of details is then in direct conflict with the desire for transparency for efficient use.

In Oberon this intrinsic dilemma has been resolved - or rather circumvented - by not including any statements for input and output at all. This extreme approach was made acceptable because of two facilities. First, there exists the module structure allowing the construction of a hierarchy of (library) modules representing various levels of abstraction. Second, Oberon permits the expression of computer specific operations, such as communication with peripheral interfaces. These operations are typically contained in modules at the lowest level of this hierarchy, and are therefore counted among the so-called *low-level facilities*. A program wishing to ignore the details of device handling imports procedures from the standard modules at the higher levels of this hierarchy. When desiring utmost efficiency or requiring access to specific properties of specific devices, one either uses low-level modules, so-called *device drivers,* or uses the primitives themselves. In the latter case, the programmer pays the price of intransportability, for his programs refer directly to particulars of either his computer or its operating system.

It is impossible to present in this context any operations of devices at the low levels of the module hierarchy, because there exists a wide variety of such devices. We restrict the following material to the presentation of two typical modules used in performing conventional input and output operations. Module *Texts* we have already encountered in examples in preceding chapters. The main concept to be presented is the *file*, data considered as *sequences* of elements of the same type. For this purpose, we present the definition of module *Files*.

We must generally distinguish between legible and illegible input and output. Legible input and output serves to communicate between the computer and its user. Mostly the data elements are characters i.e. values of type CHAR; the exception is graphical input and output. Legible data are input through keyboards, scanners, etc. Legible output is generated by displays and printers. Illegible input and output is made from and to so-called *peripheral storage media*, such as disks and tapes, but also may originate from sensors - e.g. in laboratories or drawing offices - and to devices that are

controlled by computers, such as plotters, factory assembly lines, traffic signals, and networks. Data for illegible input and output can be of any type, and need not be of type CHAR.

The vast majority of input and output operations of both the legible and illegible variety is appropriately considered as *sequential*. Their data are of a structure that does not exist as a basic data structuring method in Oberon, such as the array or the record. Sequences have the following characteristics:

1. All elements are of the same type.

2. The number of elements is not known a priori. The sequence is therefore (a simple case of) a dynamic structure. The number of elements is called the *length* of the sequence.

3. A sequence can be modified only by appending elements at its end. Appending an element is called *writing*.

4. Only a single element of a stream is visible (accessible) at any one time. Accessing this element is called *reading*, and the sequence is read by advancing from one element to the next.

In passing we state that the sequence as described above is perhaps the most successful case of data abstraction encountered. It is certainly more widely used in actual practice than the often quoted examples of stacks and queues. The language Pascal had included it among its basic data structuring methods along with arrays, records, and sets.

Before proceeding with the postulation and explanation of modules for handling sequences, which we will call *files* and *texts*, we wish to point out an important separation of function performed for legible input and output. On the one hand, there is the actual transport of data to and from the computer. This involves the activation and sensing of the state of the peripheral device, be it a keyboard, a display, or a printer. On the other hand, there is the function of transforming the representation of data. If, for example, the value of an expression of type INTEGER is to be transmitted to a display, the computer-internal representation must be transformed into the decimal representation as a sequence of digits. The display device then translates the character representation (usually consisting of 8 bits for each character) into a pattern of visible dots or lines. However, the former translation can be considered as device independent, and is therefore a prime candidate for separation from device specific operations. It can be performed by the same routines without regard whether the sequence is to be stored on a disk or to be made visible on a display.

A third class of functions that can well be separated from physical data transfer and from conversion of data representation, pertains to devices associated with more than a single sequence, the primary example being a disk store. We refer to the operations of allocating storage space and associating names with individual texts or files. Considering that texts and files are dynamic structures, storage allocation is of considerable complexity. The naming of individual files and in particular the management of directories (in order to quickly locate individual files) is another task requiring an elaborate mechanism. Both storage allocation and directory management are the tasks of individual service modules. There seem to exist as many ways to manage these tasks as there exist operating systems. And this is precisely where diversity transcends the many levels of input and output operations. We therefore strictly adhere to and confine ourselves to the abstract notion of a sequence.

22.2. Files and Riders

Elements of a sequence cannot be identified by an index, as in the case of arrays, nor by a field name, as in the case of a record. Elements are instead accessed one by one, advancing strictly sequentially. This notion itself implies the postulation of an access mechanism. In the Oberon module *Files*, we call it *Rider*. A rider is a data structure, an object, which can be placed on a file at a given position, and then be used to write or read single elements. Hence, module *Files* defines not only one, but two data types: *File* and *Rider*. The former stands for the data, the latter for the operations performed. It now follows, that several riders may be positioned on the same file, and be moved independently. Not the file has a reading or writing position, but each rider. Normally, however, one rider only is used for a file.

```
DEFINITION Files;
    TYPE File;
        Rider = RECORD eof: BOOLEAN; res: LONGINT END ;
```

```
PROCEDURE Old (name: ARRAY OF CHAR): File;
PROCEDURE New (name: ARRAY OF CHAR): File;
PROCEDURE Register (f: File);
PROCEDURE Close (f: File);
PROCEDURE Purge (f: File);
PROCEDURE Length (f: File): LONGINT;
PROCEDURE GetDate (f: File; VAR t, d: LONGINT);

PROCEDURE Set (VAR r: Rider; f: File; pos: LONGINT);
PROCEDURE Pos (VAR r: Rider): LONGINT;
PROCEDURE Base (VAR r: Rider): File;
PROCEDURE Read (VAR r: Rider; VAR ch: CHAR);
PROCEDURE ReadInt (VAR R: Rider; VAR x: INTEGER);
PROCEDURE ReadLInt (VAR R: Rider; VAR x: LONGINT);
PROCEDURE ReadReal (VAR R: Rider; VAR x: REAL);
PROCEDURE ReadLReal (VAR R: Rider; VAR x: LONGREAL);
PROCEDURE ReadString (VAR R: Rider; VAR x: ARRAY OF CHAR);

PROCEDURE Write (VAR r: Rider; ch: CHAR);
PROCEDURE WriteInt (VAR R: Rider; x: INTEGER);
PROCEDURE WriteLInt (VAR R: Rider; x: LONGINT);
PROCEDURE WriteReal (VAR R: Rider; x: REAL);
PROCEDURE WriteLReal (VAR R: Rider; x: LONGREAL);
PROCEDURE WriteString (VAR R: Rider; x: ARRAY OF CHAR);

PROCEDURE Delete (name: ARRAY OF CHAR; VAR res: INTEGER);
PROCEDURE Rename (old, new: ARRAY OF CHAR; VAR res: INTEGER);
END Files.
```

The procedures are listed in four groups. The first group operates on files, the second and the third on riders for reading and writing respectively, and the fourth on the file *directory* only. We assume that files are stored on a persistent medium, such as a magnetic disk or flash-rom, and that all files are listed in a directory with names.

In the first group, procedure *Old* yields the file listed in the directory by the specified name. *New* generates a new (empty) file with specified name. Actual registration in the directory is performed by procedure *Register*, typically after the entire file had been generated. *Close* terminates the writing of a file (flushing buffers), and it is implied in *Register*. *GetDate* yields the date and time of a file's creation.

In the second group, *Set* places a rider on a file at a specified position (between 0 and the length of the file). *Pos* indicates the rider's current position, and *Base* the file on which it is placed. The Read procedures read a specified type of data element (consisting of one or several bytes) without any transformation of representation. They advance the rider by the appropriate amount. A string is assumed to be terminated by a 0X character.

In the third group, procedures operate similarly for generating (writing) a file. Note that the rider must not necessarily be positioned at the end of the file, although this is the normal, prevalent case.

We are now in a position to show, how files are typically written and read in Oberon. Let us assume the declarations

```
f: Files.File; r: Files.Rider
```

First, the empty file is created, then a rider is associated with it, then we assume it is generated sequentially (here in a whyile-loop), and finally (and optionally) it is registered in the directory.

```
f := Files.New("MyFile"); Files.Set(r, f, 0);
WHILE more DO compute(next); Files.Write(r, next) END ;
Files.Register(f)
```

The file can thereafter be read by the following pattern. First, the file with specified name is associated with variable f. Then a rider is placed at the start of the file, and then advanced by reading.

```
f := Files.Old("MyFile"); Files.Set(r, f, 0);
WHILE ~r.eof DO Files.Read(f, next); process(next) END
```

Note that the rider indicates that the *end of the file* had been reached after the first unsuccessful attempt of reading.

22.3. Texts, Readers and Writers

Texts, basically, are sequences of characters. In the Oberon System, however, texts have some additional properties. As they can be displayed and printed, they must carry properties which determine their style and appearance. In particular, a *font* is specified. This is a style, the patterns by which characters are represented. Subsequences of charcaters in a text may have their individual fonts. Furthermore, Oberon texts specify color and vertical offset (allowing for negative offset for subscripting, and positive offset for superscripting). Texts are typically subjected to complicated editing operations, which require a flexible, internal data representation. Therefore Texts in Oberon differ substantiall from Files in many respect. However, in their essence they are also sequences, and the basic operations strongly resemble those of files.

Our special interest in texts is justified by the fact, that humans communicate with computers mostly via texts, be they typed on a keyboard or displayed on a screen. As a consequence, reading and writing of texts usually includes a change of data representation. For example, integers will have to be changed into sequences of decimal digits on output, and the reading of integers requires the reading of a sequence of digits and the computation of the represented value. These conversions are included in the read/write procedures of module *Texts*, which we have used in many examples in preceding chapters. The following is its definition, which evidently resembles that of *Files*. In place of the single type *Rider* we find the two types *Reader* and *Writer*.

```
DEFINITION Texts;
    IMPORT Files, Fonts;
    CONST replace = 0; insert = 1; delete = 2;
          Inval = 0; Name = 1; String = 2; Int = 3; Real = 4; LongReal = 5; Char = 6;
    TYPE Text = POINTER TO RECORD
                    len: LONGINT
              END ;
          Buffer;
          Reader = RECORD
                    eot: BOOLEAN;
                    fnt: Fonts.Font;
                    col, voff: SHORTINT
              END ;

          Scanner = RECORD (Reader)
                    nextCh: CHAR;
                    line, class: INTEGER;
                    i: LONGINT;
                    x: REAL;
                    y: LONGREAL;
                    c: CHAR;
                    len: SHORTINT;
                    s: ARRAY 32 OF CHAR;
              END ;

          Writer = RECORD
                    buf: Buffer;
                    fnt: Fonts.Font;
                    col, voff: SHORTINT;
              END ;
    PROCEDURE Open(t: Text; name: ARRAY OF CHAR);
    PROCEDURE Delete(t: Text; beg, end: LONGINT);
    PROCEDURE Insert(t: Text; pos: LONGINT; b: Buffer);
    PROCEDURE Append(t: Text; b: Buffer);
    PROCEDURE ChangeLooks
        (T: Text; beg, end: LONGINT; sel: SET; fnt: Fonts.Font; col, voff: SHORTINT);
```

```
PROCEDURE OpenReader (VAR r: Reader; t: Text; pos: LONGINT);
PROCEDURE Read (VAR r: Reader; VAR ch: CHAR);
PROCEDURE Pos (VAR r: Reader): LONGINT;
PROCEDURE OpenScanner (VAR S: Scanner; t: Text; pos: LONGINT);
PROCEDURE Scan (VAR S: Scanner);

PROCEDURE OpenWriter (VAR w: Writer);
PROCEDURE SetFont (VAR w: Writer; fnt: Fonts.Font);
PROCEDURE SetColor (VAR w: Writer; col: SHORTINT);
PROCEDURE SetOffset (VAR w: Writer; voff: SHORTINT);
PROCEDURE Write (VAR w: Writer; ch: CHAR);
PROCEDURE WriteLn (VAR w: Writer);
PROCEDURE WriteString (VAR w: Writer; s: ARRAY OF CHAR);
PROCEDURE WriteInt (VAR w: Writer; x, n: LONGINT);
PROCEDURE WriteHex (VAR w: Writer; x: LONGINT);
PROCEDURE WriteReal (VAR w: Writer; x: REAL; n: INTEGER);
PROCEDURE WriteRealFix (VAR w: Writer; x: REAL; n, k: INTEGER);
PROCEDURE WriteLongReal (VAR w: Writer; x: LONGREAL; n: INTEGER);
PROCEDURE WriteDate (VAR w: Writer; t, d: LONGINT);

PROCEDURE Load (VAR r: Files.Rider; t: Text);
PROCEDURE Store (VAR r: Files.Rider; t: Text)
END Texts.
```

Procedure *OpenWriter* uses as defaults a standard font, black color, and zero offset.

Simple, sequential reading and writing of a text now follow the patterns of reading and writing a file. Let us assume the following declarations of variables:

```
T: Texts.Text;
R: Texts.Reader;
W: Texts.Writer;

NEW(T); Texts.Open(T, "MyName"); Texts.OpenReader(R, T, 0); Texts.Read(R, ch);
WHILE ~R.eot DO process(ch); Texts.Read(R, ch) END

NEW(T); Texts.Open(T, "MyName"); Texts.OpenWriter(W); generate(ch)
WHILE more DO Texts.Write(W, ch); generate(ch) END ;
Texts.Append(T, W.buf)
```

Note that the Oberon paradigm is to write a text, or a piece of text, first into a buffer, and thereafter insert or append it to a text. This is done for reasons of efficiency, because the possibly needed rendering of the text, for example on a display, can be done once upon insertion of the buffered piece of text rather than after generating each character.

Very often one wants to read a text consisting of a sequence of items which are not all of the same type, such as numbers, strings, names, etc. When using procedures reading items of a fixed type, the programmer must know the exact sequence in which the various items will appear in the text. Typically one does not know, and even if a specific order is specified, mistakes may occur. So one should like to have available a reading mechanism that reads items one at a time, but lets the program determine what type of item was read, and take further steps accordingly. This facility is provided in the Oberon text module by the machanism called *scanning*. In place of a *Reader* we use a *Scanner*. Procedure *Scan(S)* then reads the next item. Its kind can be determined by inspecting the field *S.class*, and its value accordingly is given by *S.i* in the case an integer was read, *S.x* if a real number was read, *S.s* if a name or a string was read. This scheme has proven to be most useful because of its flexibility. It defines the following syntax for texts to be scanned. Blanks and line ends serve to separate consecutive items.

```
item     =     name | integer | real | longreal | string | SpecialChar.
name     =     letter {letter | digit}.
integer  =     [sign] digit {digit} | digit {digit | hexdig} „H".
hexdig   =     „A" | „B" | „C" | „D" | „E" | „F".
sign     =     [„-"].
real     =     [sign] digit {digit} „." digit {digit} [„E" [sign] digit {digit}].
```

longreal   =   [sign] digit {digit} „.“ digit {digit} „D“ [sign] digit {digit}.

A string is any sequence of any characters (except quotes) enclosed in quotes. Special characters are are all characters except letters, digits, and the quote mark.

22.4. Standard Input and Output

In order to simplify the description of input and output of texts in simple cases, the Oberon system introduces some conventions and global variables. We call this source of input and sink of output *standard* input and output.

Standard output sink is the text *Log* defined as global variable in module *Oberon*. Given a global writer W, the text, previously written by Write procedures into the writer's buffer (W.buf) is sent to the displayed *Log* text by the statement

Texts.Append(Oberon.Log, W.buf)

The standard input is often assumed to be the text following the command which invoked execution of a given procedure. This input text is identified by the variable *Par* in module *Oberon*. It is considered as parameter of the invoked command, and it is typically read by the scanning mechanism described above. The necessary statements are:

Texts.OpenScanner(S, Oberon.Par.text, Oberon.Par.pos); Texts.Scan(S)

The first item thus read may be the first item of the desired text, or it may be the name of the text file to be scanned, or something else, according to the specification of the individual command. The following conventions have established themselves for the designation of input texts. Assume that the command (procedure) name is followed by

- an identifier (possibly qualified). Then this is the name of the input text (input file),
- an asterisk (*). Then the marked viewer (window) contains the input text,
- an @ symbol. Then the most recent text selection is the beginning of the input text,
- an arrow (^). Then the most recent selection denotes the file name of the input text.

These conventions are expressed by the following function procedure yielding the designated input text:

```
PROCEDURE This*(): Texts.Text;
  VAR beg, end, time: LONGINT;
    S: Texts.Scanner; T: Texts.Text; v: Viewers.Viewer;
  BEGIN Texts.OpenScanner(S, Oberon.Par.text, Oberon.Par.pos); Texts.Scan(S);
    IF S.class = Texts.Char THEN
      IF S.c = "*" THEN  (*input in marked viewer*)
        v := Oberon.MarkedViewer();
        IF (v.dsc # NIL) & (v.dsc.next IS TextFrames.Frame) THEN
          T := v.dsc.next(TextFrames.Frame).text; beg := 0
        END
      ELSIF S.c = "@" THEN  (*input starts at selection*)
        Oberon.GetSelection(T, beg, end, time);
        IF time < 0 THEN T := NIL END
      ELSIF S.c = "^" THEN  (*selection is the file name*)
        Oberon.GetSelection(T, beg, end, time);
        IF time >= 0 THEN (*there is a selection*)
          Texts.OpenScanner(S, Oberon.Par.text, Oberon.Par.pos); Texts.Scan(S);
          IF S.class = Texts.Name THEN  (*input is named file*)
            NEW(T); Texts.Open(T, S.s); pos := 0
          END
        END
      END
    ELSIF S.class = Texts.Name THEN  (*input is named file*)
      NEW(T); Texts.Open(T, S.s); pos := 0
    END ;
    RETURN T
  END This;
```

Let us assume that this function procedure is defined in module *Oberon*, together with the global variable *pos*, indicating the position in the text where input is to start. The following program for copying a text may serve as a pattern for selecting the input according to the conventions postulated above.

```
MODULE ProgramPattern;
  IMPORT Texts, Oberon;
  VAR W: Texts.Writer;  (*global writer*)

  PROCEDURE copy(VAR R: Texts.Reader);
    VAR ch: CHAR;
  BEGIN Texts.Read(R, ch);
    WHILE ~R.eot DO Texts.Write(W, ch); Texts.Read(R, ch) END
  END copy

  PROCEDURE SomeCommand*;
    VAR R: Texts.Reader;  (*local reader*)
  BEGIN Texts.OpenReader(R, Oberon.This.text, Oberon.Par.pos);
    copy(R, W); Texts.Append(Oberon.Log, W.buf)
  END SomeCommand;

BEGIN Texts.OpenWriter(W)
END ProgramPattern.
```

Module *Oberon* can be considered as the core of the Oberon system housing a small number of global resources. These include the previously encountered output *Log*, and the record *Par*, specifying the viewer, frame, and text in which the activated command lies, and hence providing access to its input parameters. Here we show only an excerpt of the definition of module Oberon:

```
TYPE ParList = POINTER TO RECORD
            vwr: Viewers.Viewer;
            frame: Display.Frame;
            text: Texts.Text;
            pos: LONGINT;
    END ;

VAR Log: Texts.Text;
  Par: ParList;

PROCEDURE Time (): LONGINT;

PROCEDURE AllocateUserViewer (DX: INTEGER; VAR X, Y: INTEGER);
PROCEDURE AllocateSystemViewer (DX: INTEGER; VAR X, Y: INTEGER);
  (*provide coordinates X and Y for a new viewer to be allocated*)

PROCEDURE GetSelection (VAR text: Texts.Text; VAR beg, end, time: LONGINT);
```

To conclude this introduction to Oberon conventions about input and output, we show how a new viewer (window) is opened, and how the output text is directed into this text viewer. Optimal positioning of the new viewer is achieved by procedure *Oberon.AllocateViewer*, with the specified set of standard, frequently needed commands in the title bar.

```
PROCEDURE OpenViewer(T: Texts.Text);
  VAR V: MenuViewers.Viewer; X, Y: INTEGER;
BEGIN Oberon.AllocateUserViewer(Oberon.Mouse.X, X, Y);
  V := MenuViewers.New(TextFrames.NewMenu(
          "Name", "System.Close  System.Copy  System.Grow  Edit.Search  Edit.Store"),
          TextFrames.NewText(T, 0), TextFrames.menuH, X, Y)
END OpenViewer;
```

Appendix

1. The Syntax of Oberon

ident = letter {letter | digit}.
number = integer | real.
integer = digit {digit} | digit {hexDigit} "H" .
real = digit {digit} "." {digit} [ScaleFactor].
ScaleFactor = ("E" | "D") ["+" | "-"] digit {digit}.
hexDigit = digit | "A" | "B" | "C" | "D" | "E" | "F".
digit = "0" | "1" | "2" | "3" | "4" | "5" | "6" | "7" | "8" | "9".
CharConstant = """ character """ | digit {hexDigit} "X".
string = """ {character} """ .
qualident = [ident "."] ident.
identdef = ident ["*"].

TypeDeclaration = identdef "=" type.
type = qualident | ArrayType | RecordType | PointerType | ProcedureType.
ArrayType = ARRAY length {"," length} OF type.
length = ConstExpression.
RecordType = RECORD ["(" BaseType ")"] FieldListSequence END.
BaseType = qualident.
FieldListSequence = FieldList {";" FieldList}.
FieldList = [IdentList ":" type].
IdentList = identdef {"," identdef}.
PointerType = POINTER TO type.
ProcedureType = PROCEDURE [FormalParameters].

VariableDeclaration = IdentList ":" type.

designator = qualident {"." ident | "[" ExpList "]" | "(" qualident ")" | "^" }.
ExpList = expression {"," expression}.
expression = SimpleExpression [relation SimpleExpression].
relation = "=" | "#" | "<" | "<=" | ">" | ">=" | IN | IS.
SimpleExpression = ["+"|"-"] term {AddOperator term}.
AddOperator = "+" | "-" | OR .
term = factor {MulOperator factor}.
MulOperator = "*" | "/" | DIV | MOD | "&" .
factor = number | CharConstant | string | NIL | set |
        designator [ActualParameters] | "(" expression ")" | "~" factor.
set = "{" [element {"," element}] "}".
element = expression [".." expression].
ActualParameters = "(" [ExpList] ")" .

statement = [assignment | ProcedureCall |
        IfStatement | CaseStatement | WhileStatement | RepeatStatement |
        LoopStatement | WithStatement | EXIT | RETURN [expression]].
assignment = designator ":=" expression.
ProcedureCall = designator [ActualParameters].
IfStatement = IF expression THEN StatementSequence
        {ELSIF expression THEN StatementSequence}
        [ELSE StatementSequence]
        END.
CaseStatement = CASE expression OF case {"|" case} [ELSE StatementSequence] END.
Case = [CaseLabelList ":" StatementSequence].
CaseLabelList = CaseLabels {"," CaseLabels}.
CaseLabels = ConstExpression [".." ConstExpression].
WhileStatement = WHILE expression DO StatementSequence END.
LoopStatement = LOOP StatementSequence END.
WithStatement = WITH qualident ":" qualident DO StatementSequence END .

ProcedureDeclaration = ProcedureHeading ";" ProcedureBody ident.

ProcedureHeading = PROCEDURE identdef [FormalParameters].
ProcedureBody = DeclarationSequence [BEGIN StatementSequence] END.
ForwardDeclaration = PROCEDURE "^" identdef [FormalParameters].
FormalParameters = "(" [FPSection {";" FPSection}] ")" [":" qualident].
FPSection = [VAR] ident {"," ident} ":" FormalType.
FormalType = {ARRAY OF} qualident.

DeclarationSequence = {CONST {ConstantDeclaration ";"} |
      TYPE {TypeDeclaration ";"} | VAR {VariableDeclaration ";"}}
      {ProcedureDeclaration ";" | ForwardDeclaration ";"}.
Module = MODULE ident ";" [ImportList] DeclarationSequence
      [BEGIN StatementSequence] END ident "." .
ImportList = IMPORT import {"," import} ";" .
Import = ident [":=" ident].

2. Symbols and Keywords

+	:=	ARRAY	IS	TO
-	^	BEGIN	LOOP	TYPE
*	=	CASE	MOD	UNTIL
/	#	CONST	MODULE	VAR
~	<	DIV	NIL	WHILE
&	>	DO	OF	WITH
.	<=	ELSE	OR	
,	>=	ELSIF	POINTER	
;	..	END	PROCEDURE	
		:	EXIT	RECORD	
(	)	IF	REPEAT	
[	]	IMPORT	RETURN	
{	}	IN	THEN	

3. Standard Functions and Procedures

Name	Argument type	Result type	
ABS(x)	numeric	type of x	absolute value
ODD(x)	integer type	BOOLEAN	x MOD 2 = 1
CAP(x)	CHAR	CHAR	corresponding capital letter
ASH(x, n)	integer types	LONGINT	$x \times 2^n$ (arithmetic shift)
LEN(v, n)	v: array type	LONGINT	length of v in dimension n
LEN(v)	v: array type	LONGINT	length of v in dimension 0
ORD(x)	CHAR	INTEGER	ordinal number of character x
CHR(x)	INTEGER	CHAR	character with ordinal number x
SHORT(x)	LONGINT	INTEGER	identity (truncation possible!)
	INTEGER	SHORTINT	
	LONGREAL	REAL	
LONG(x)	SHORTINT	INTEGER	identity
	INTEGER	LONGINT	
	REAL	LONGREAL	
ENTIER(x)	real type	LONGINT	largest integer not graeter x
INC(v, n)	integer types		v := v + n
INC(v)	integer types		v := v + 1
DEC(v, n)	integer types		v := v − n
DEC(v)	integer types		v := v -1
INCL(v, n)	v: SET; n: integer type		v := v + {n}
EXCL(v, n)	v: SET; n: integer type		v := v - {n}
COPY(x, v)	x: character arry, string		v := x
NEW(v)	pointer type		allocate v^
HALT(x)	integer constant		terminate computation